



Introduction to Transformers



Michel RIVEILL
michel.riveill@univ-cotedazur.fr

General word/sentences representations

- ▶ **Feature-based approaches (or static approach)**
 - ▶ Non-neural word representations (BOW)
 - ▶ Neural embedding
 - ▶ Word embedding:
 - Word2Vec, Glove,...
 - ▶ Sentence embedding or Paragraph embedding,...
- ▶ **Embeddings from Language Models**
 - ▶ Replace static embeddings (lexicon lookup) with context-dependent embeddings (produced by a deep neural language model)
 - ▶ Each token's representation is a function of the entire input sentence, computed by a deep (multi-layer) bidirectional language model
 - ▶ Return for each token a (task-dependent) linear combination of its representation across layers.
 - ▶ Different layers capture different information

Embeddings from Language Models

- ▶ Deep contextualised word representation
 - ▶ **ELMo**, Embeddings from Language Models, Peters et al., 2018
- ▶ Fine-tuning approaches
 - ▶ **GPT**
 - ▶ Generative Pre-trained Transformer, Radford et al., 2018
 - ▶ **BERT**
 - ▶ Bi-directional Encoder Representations from Transformers, Devlin et al., 2018




Elmo



ELMo: deep contextualised word representation (Peters et al., 2018)

- ▶ “Instead of using a fixed embedding for each word, ELMo looks at the entire sentence before assigning each word in it an embedding.”

ELMo
Embeddings



Words to embed

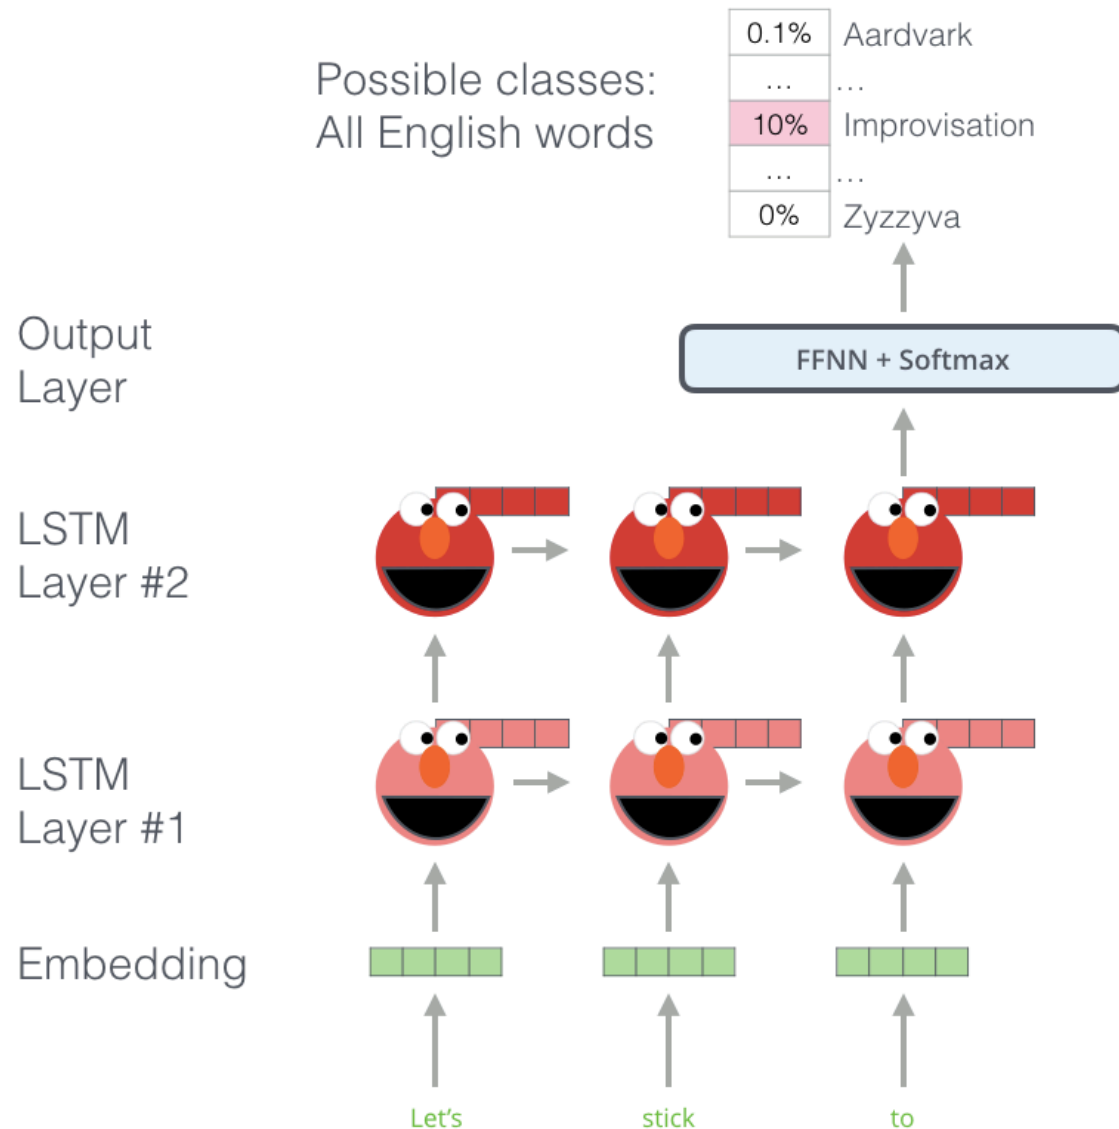


ELMo's secret

ELMo was trained in an unsupervised manner, like word2vec:

- predicting the next word in a sequence of words - a task called "language modeling".

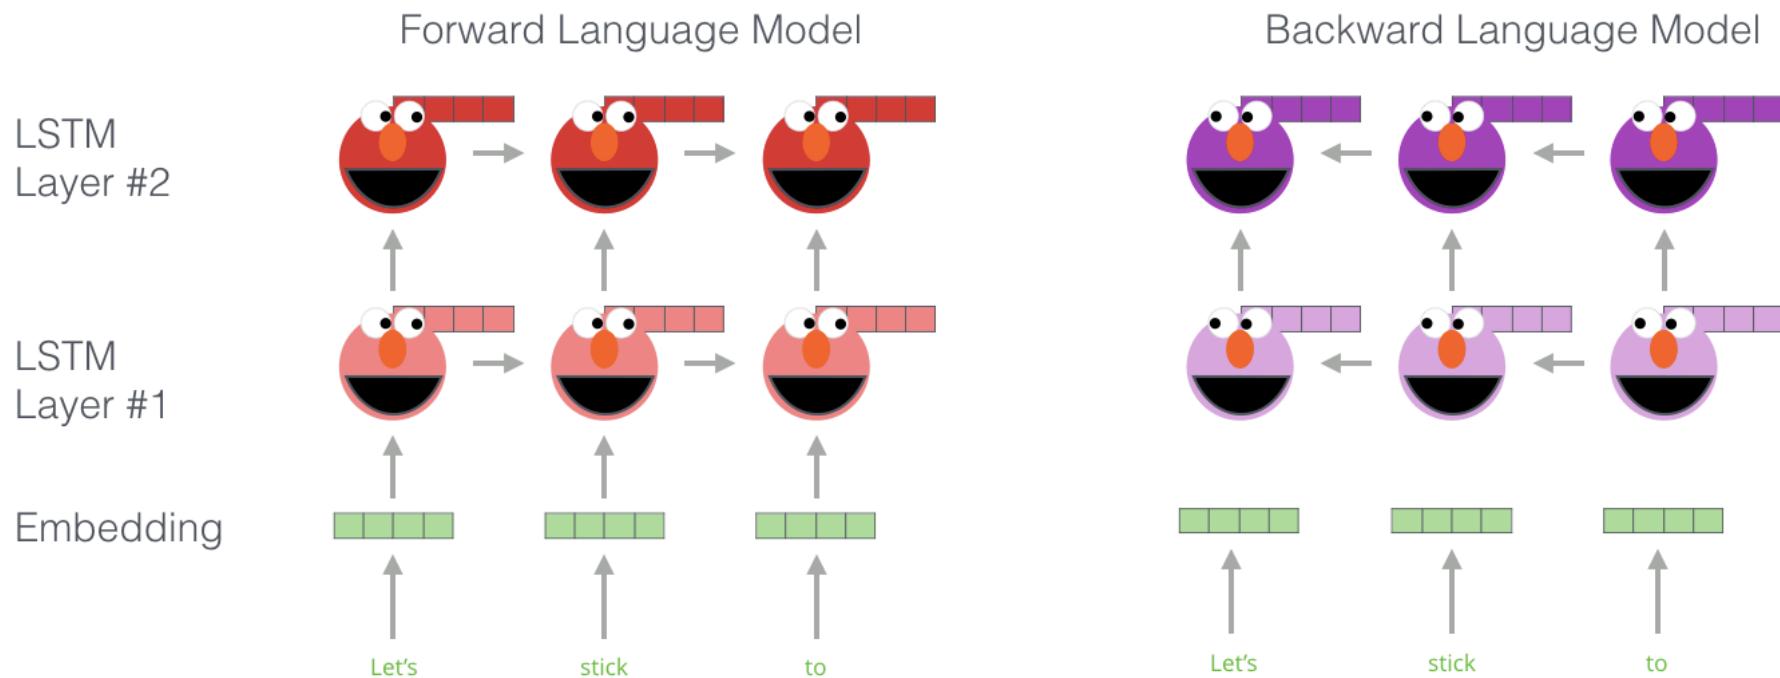
This is useful because there is a large amount of unlabeled text data available.



ELMo's secret

- In practice ELMo uses a bi-directional MSTL.
- His linguistic model tries to capture the relationships within a sentence in both directions.

Embedding of “stick” in “Let’s stick to” - Step #1




ELMo's secret

ELMo proposes a contextualized embedding by grouping hidden states (and initial integration) in a certain way (concatenation followed by a weighted summation).

Embedding of “stick” in “Let’s stick to” - Step #2

1- Concatenate hidden layers



2- Multiply each vector by a weight based on the task

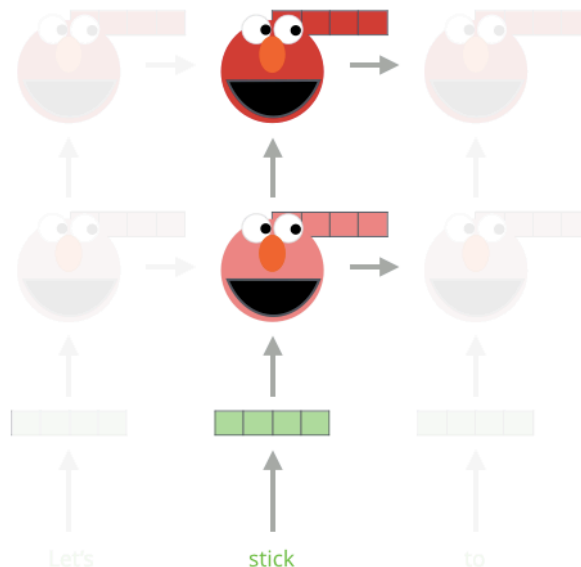


3- Sum the (now weighted) vectors




ELMo embedding of “stick” for this task in this context

Forward Language Model



Backward Language Model



ELMO architecture

- ▶ **Input representation**
 - ▶ input token representations are purely character-based: a character CNN, followed by linear projection to reduce dimensionality
 - ▶ “2048 characters n-gram convolutional filters with two highway layers, followed by a linear projection to 512 dimensions”
 - ▶ Advantage over using fixed embeddings: no UNK tokens, any word can be represented
- ▶ **Train a multi-layer bidirectional language model with character convolutions on raw text**
 - ▶ The forward LM is a deep LSTM that goes over the sequence from start to end to predict token t_k based on the prefix $t_1 \dots t_{k-1}$
 - ▶ The backward LM is a deep LSTM that goes over the sequence from end to start to predict token t_k based on the suffix $t_{k+1} \dots t_N$
- ▶ **Each layer of this language model network computes a vector representation for each token**
 - ▶ Train these LMs jointly, with the same parameters for the token representations and the softmax layer (but not for the LSTMs)
- ▶ **Freeze the parameters of the language model**
- ▶ **For each task: train task-dependent softmax weights to combine**
 - ▶ the layer-wise representations into a single vector for each token
 - ▶ *jointly* with a task-specific model that uses those vectors

How ELMo different from other word embeddings ?

- ▶ Suppose we have a couple of sentences:
 - ▶ I **read** the book yesterday.
 - ▶ Can you **read** the letter now?
- ▶ **Polysemy** : a word have multiple meanings or senses
 - ▶ “read” in the first sentence is in the past tense.
 - ▶ “read” in the second sentence is in the present tense
- ▶ Embedding of word read
 - ▶ With Keras embedding, word2vec, glove:
 - ▶ Read have always the same embedding
 - ▶ With ELMo:
 - ▶ Read have a contextualized embedding

What we can do with ELMo?

- ▶ ELMo allows to build an embedding for a list of sentences
 - ▶ It is then possible to couple this embedding with another LogisticRegression model, MLP, etc. for sentiment analysis tasks.
- ▶ It is of course possible to couple ELMo with other neural network-based models to perform more complex tasks.
 - ▶ Machine Translation
 - ▶ Language Modeling
 - ▶ Text Summarization
 - ▶ Named Entity Recognition
 - ▶ Question-Answering Systems
- ▶ It is possible to use embedding as or on the contrary, decide to make fine grained tuning

How to embed sentences with ELMo

<https://allennlp.org/elmo>

- ▶ See notebook ELMO
 - ▶ Need the installation of tensorflow_hub
- ▶ `elmo = hub.KerasLayer("https://tfhub.dev/google/elmo/2", trainable=False)`
- ▶ `text = ["A long sentence.",
 "Single-word",
 "URL http://example.com"]`
- ▶ `embeddings = elmo(tf.convert_to_tensor(np.asarray(text)))`
- ▶ `embeddings.shape`
`TensorShape([3, 1024])`
- ▶ Reuse the embedding for another task



Transformers



Transformers

Originally


- ▶ Sequence transduction model based on attention
 - ▶ no convolutions or recurrence
 - ▶ easier to parallelize than recurrent nets
 - ▶ faster to train than recurrent nets
 - ▶ captures more long-range dependencies than CNNs (Convolutional Neural Network) with fewer parameters

Now

- ▶ Transformers use stacked self-attention and pointwise, fully-connected layers for the encoder and decoder

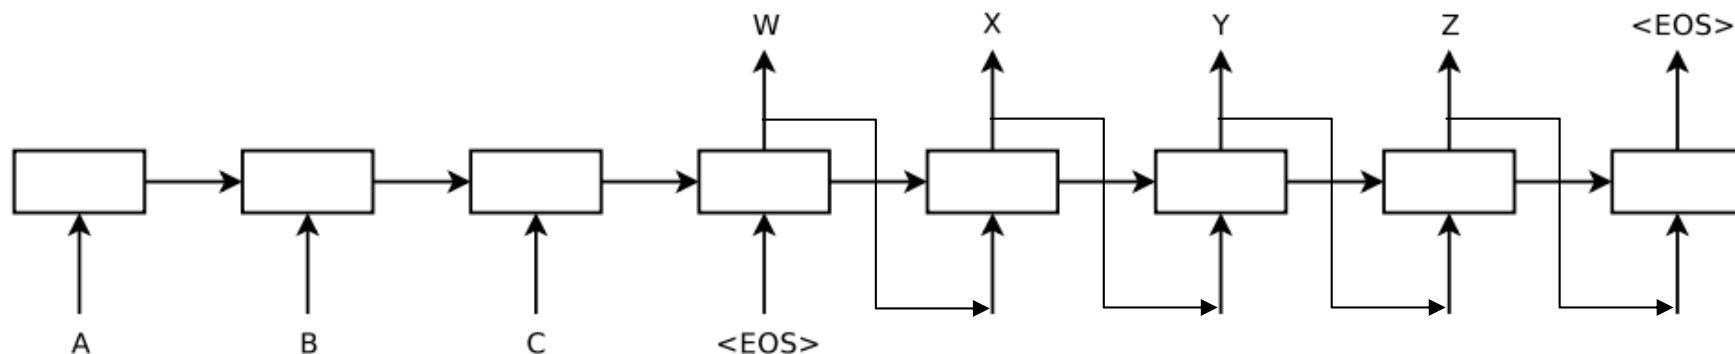
Seq2Seq architecture and Attention mechanism

- Brings many advances in NLP tasks




Seq2seq architecture


- ▶ Seq2Seq is a two-part deep learning architecture to map sequence inputs into sequence outputs
 - ▶ was initially proposed for the machine translation task
 - ▶ but can be applied for other sequence-to-sequence mapping
- ▶ Built using two Recurrent Neural Networks (RNNs), namely the encoder and the decoder
 - ▶ The encoder reads a sequence input with variable lengths, e.g., English words,
 - ▶ and the decoder produces a sequence output, e.g., corresponding French words, considering the hidden state from the encoder. The hidden state
- ▶ Main problem: sends source information from the encoder to the decoder, linking the two. Both the encoder and decoder consist of RNN cells or its variants such as LSTM and GRU.
 - ▶ difficult to parallelize, very long learning time



Replace Sequence by Self-attention




Difficult to parallelize
with RNNs




Idea: Replace RNN with
CNNs (CNN can parallel)

Replace Sequence by Self-attention

b^1, b^2, b^3, b^4 can be
parallelly computed.



Difficult to parallelize
with RNNs




Idea: Replace RNN with
CNNs (CNN can parallel)

Replace Sequence by Self-attention

b^i is obtained based on the whole input sequence.

b^1, b^2, b^3, b^4 can be parallelly computed.




You can try to replace any thing that has been done by RNN with self-attention.

LSTM Encoder-decoder machine

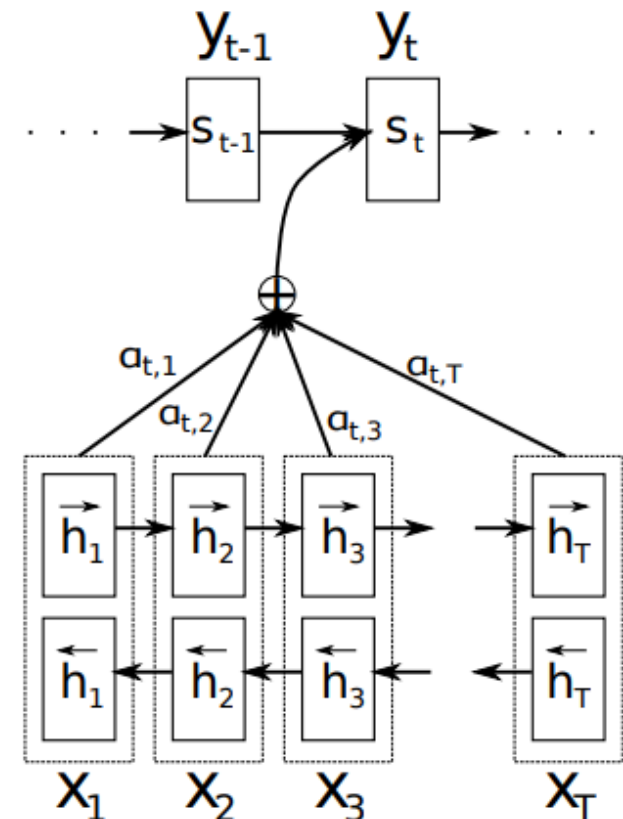
Seq2Seq
architecture cannot
capture all
information by a
single fixed length
vector (i.e. the
hidden state of the
encoder)

Problems when
processing long
sequences
(vanishing gradient
Problem)




Attention

- ▶ Attention proposes to use a context vector to represent the contributions of the source and the target (s_t)
- ▶ Context vector preserves the information of all hidden states of the encoder cells and aligns them with the current target output ($\sum_i \alpha_{t,i} \cdot h_i$)
 - ▶ $\sum_i \alpha_{t,i} = 1$
- ▶ Model to "take care" of a certain part of the source inputs
 - ▶ $\alpha_{t,i} \approx 1, \text{important}$
 - ▶ $\alpha_{t,i} \approx 0, \text{not important}$



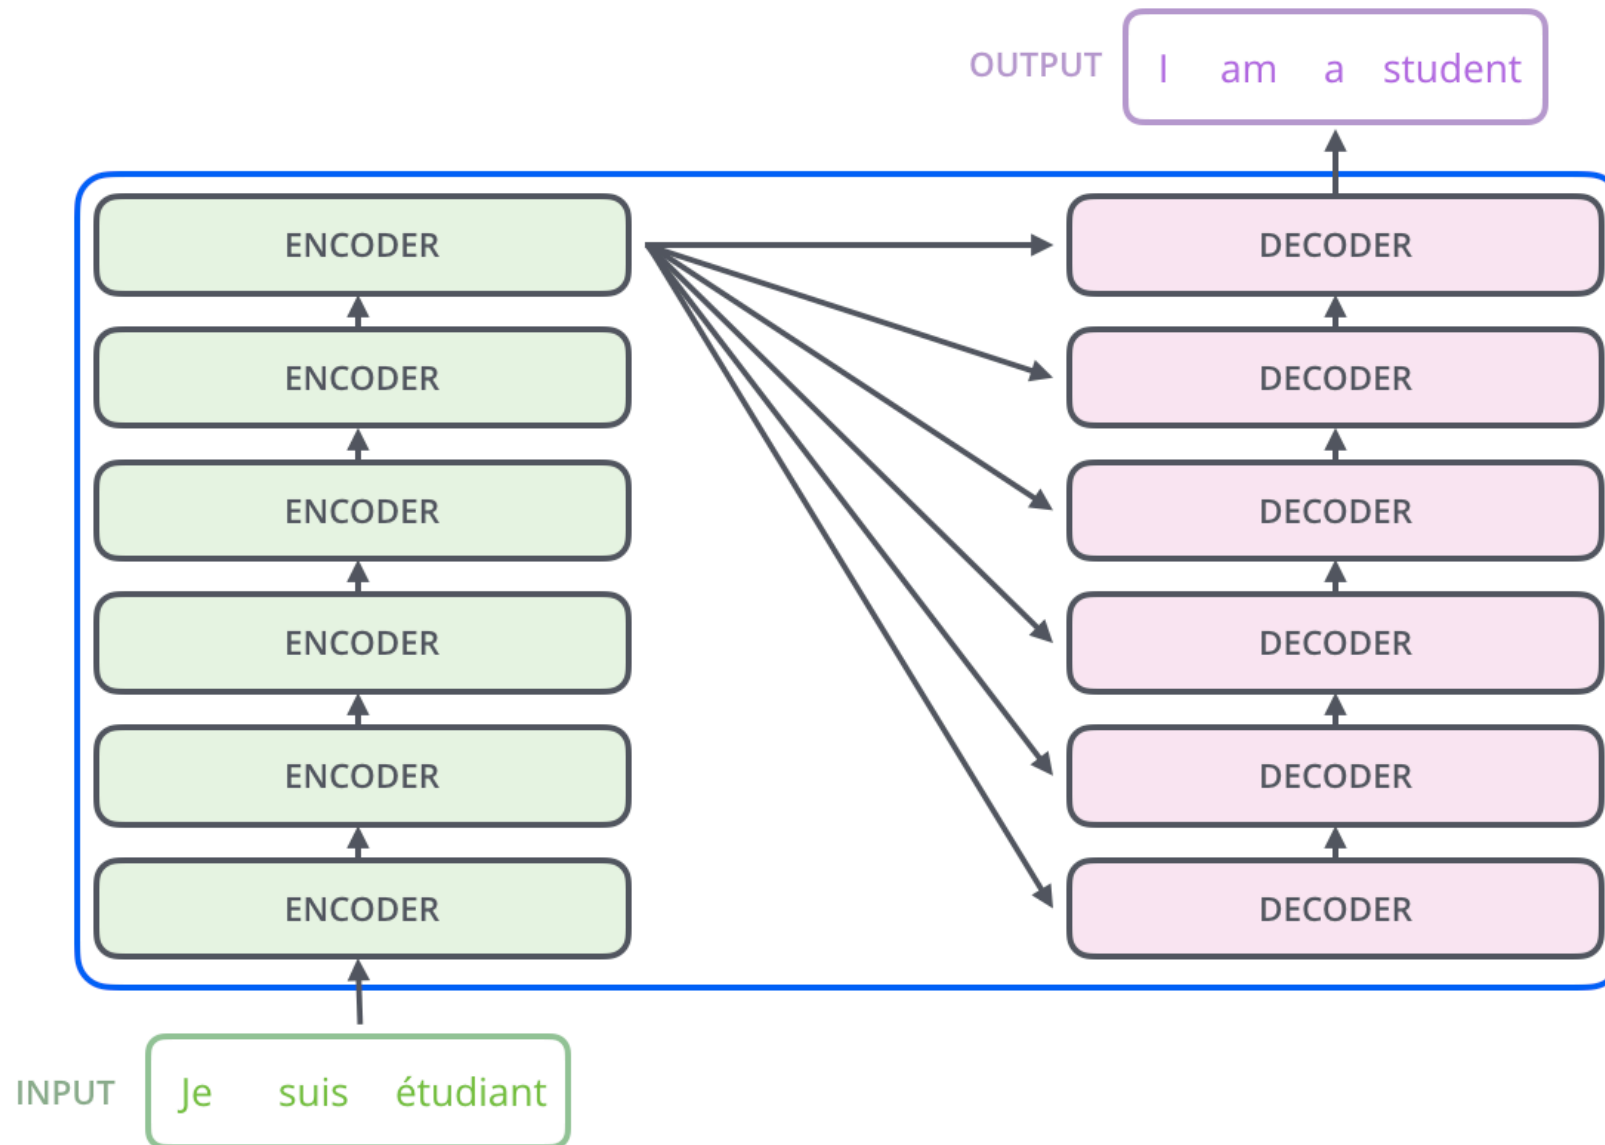
Attention



Transformers

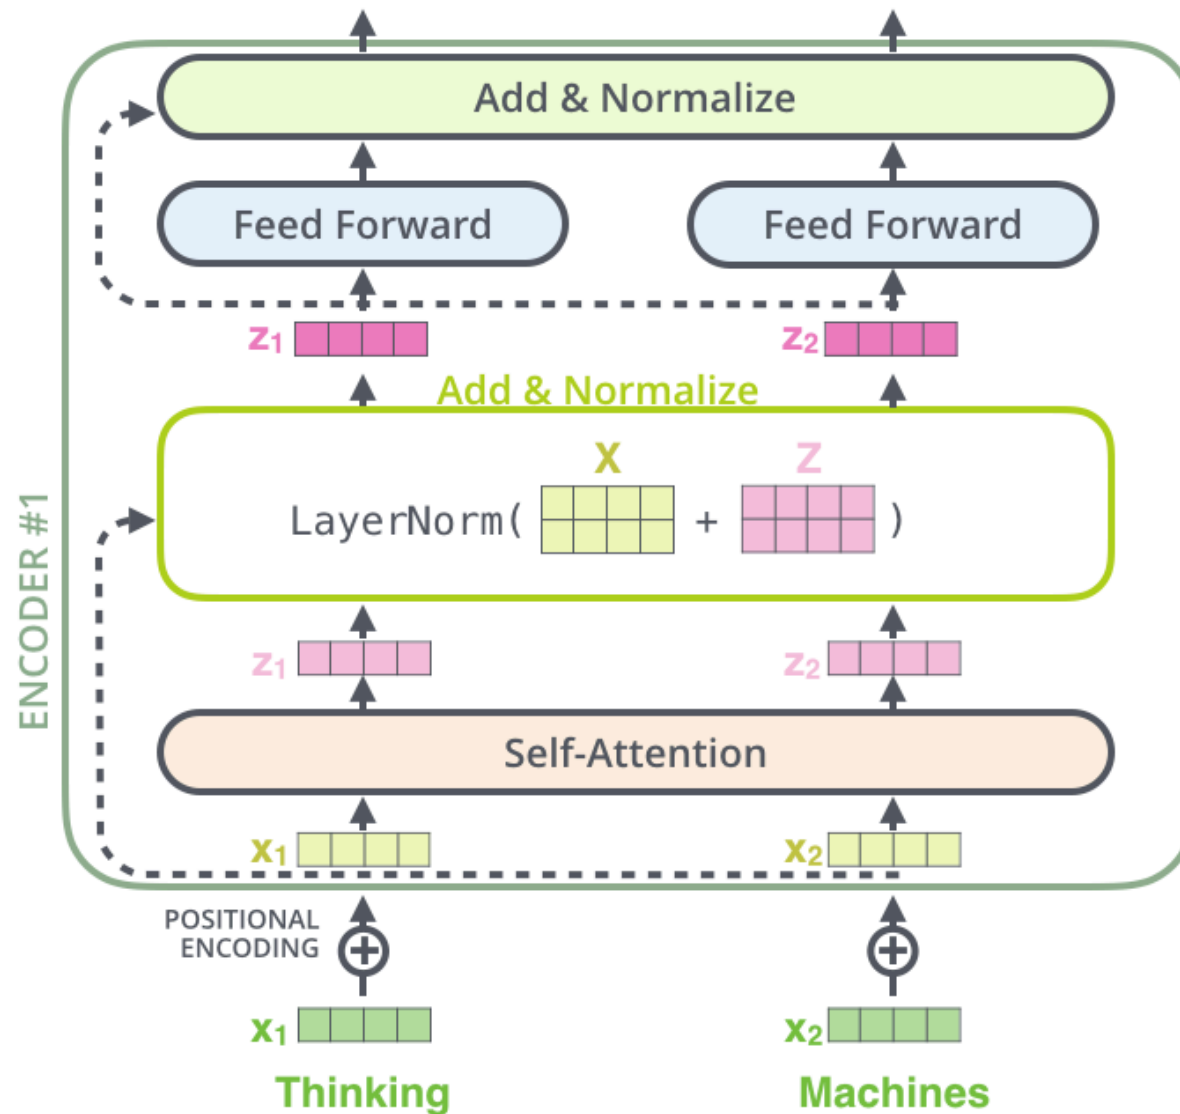
- ▶ Models that use
 - ▶ Extend Seq2seq architecture with self attention
 - ▶ Use adapted self-attention to LSTM network
- ▶ Two main blocks
 - ▶ Encoder
 - ▶ Decoder

Transformer: Going beyond LSTMs



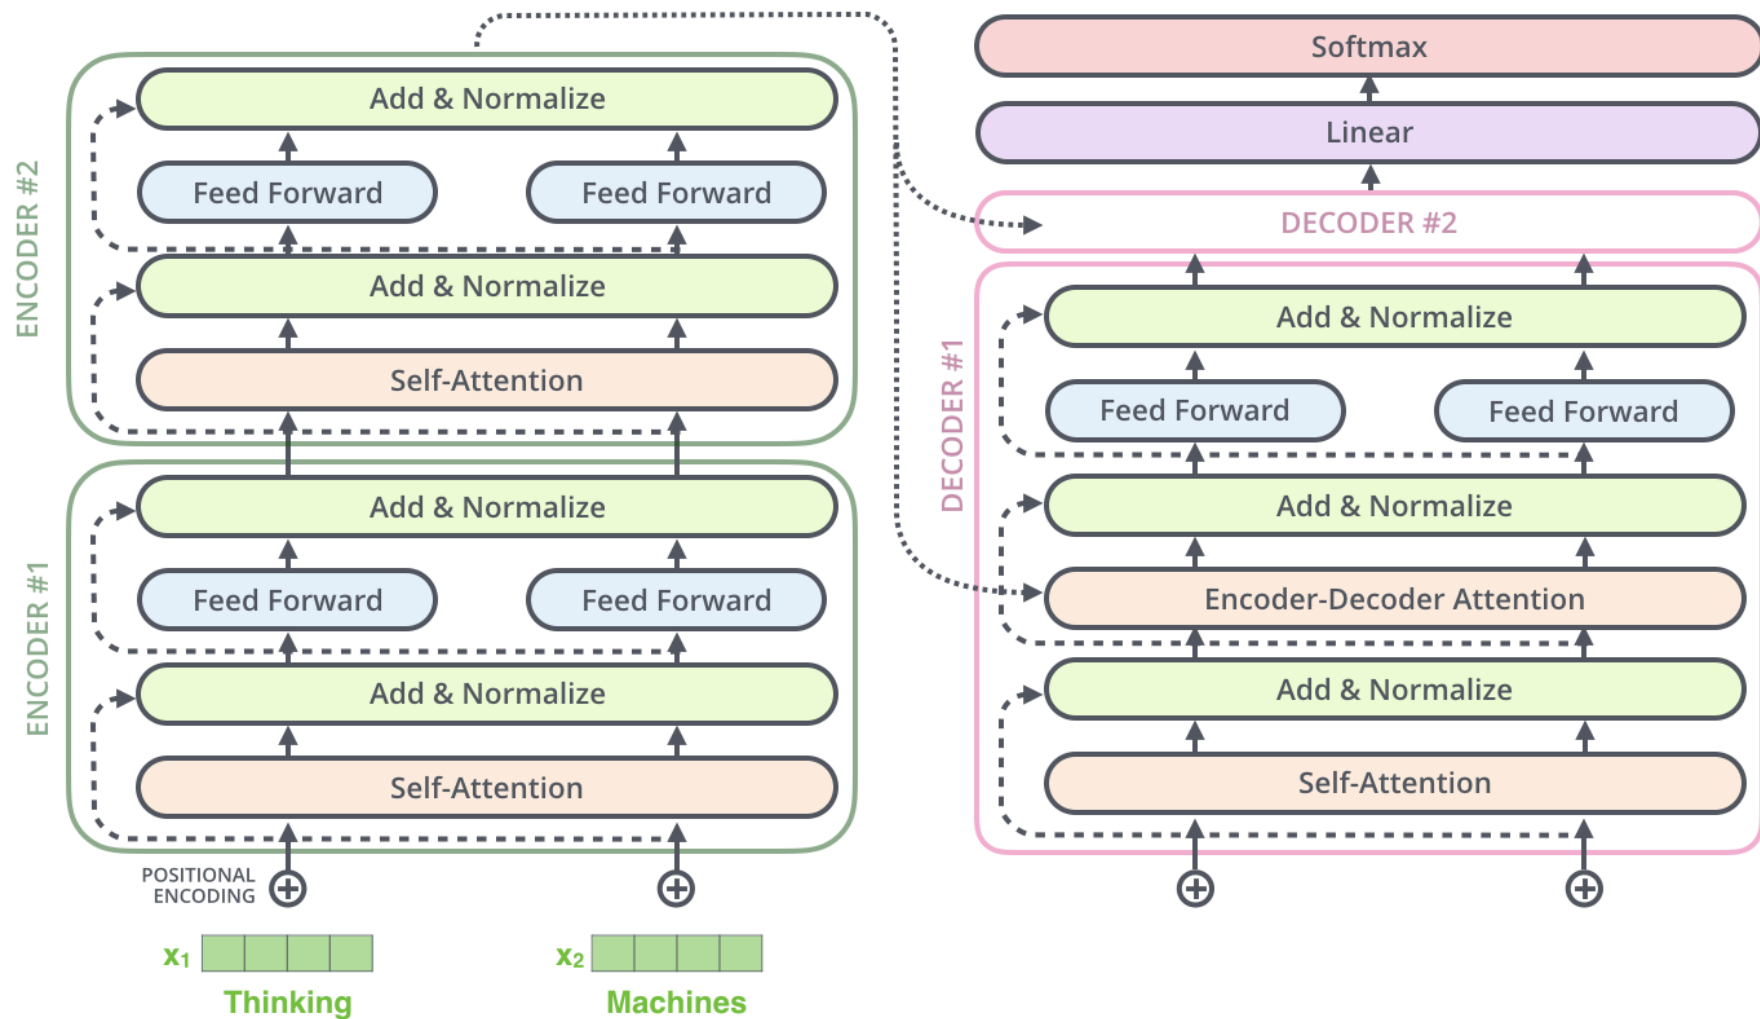
Transformer

The encoder



Transformer

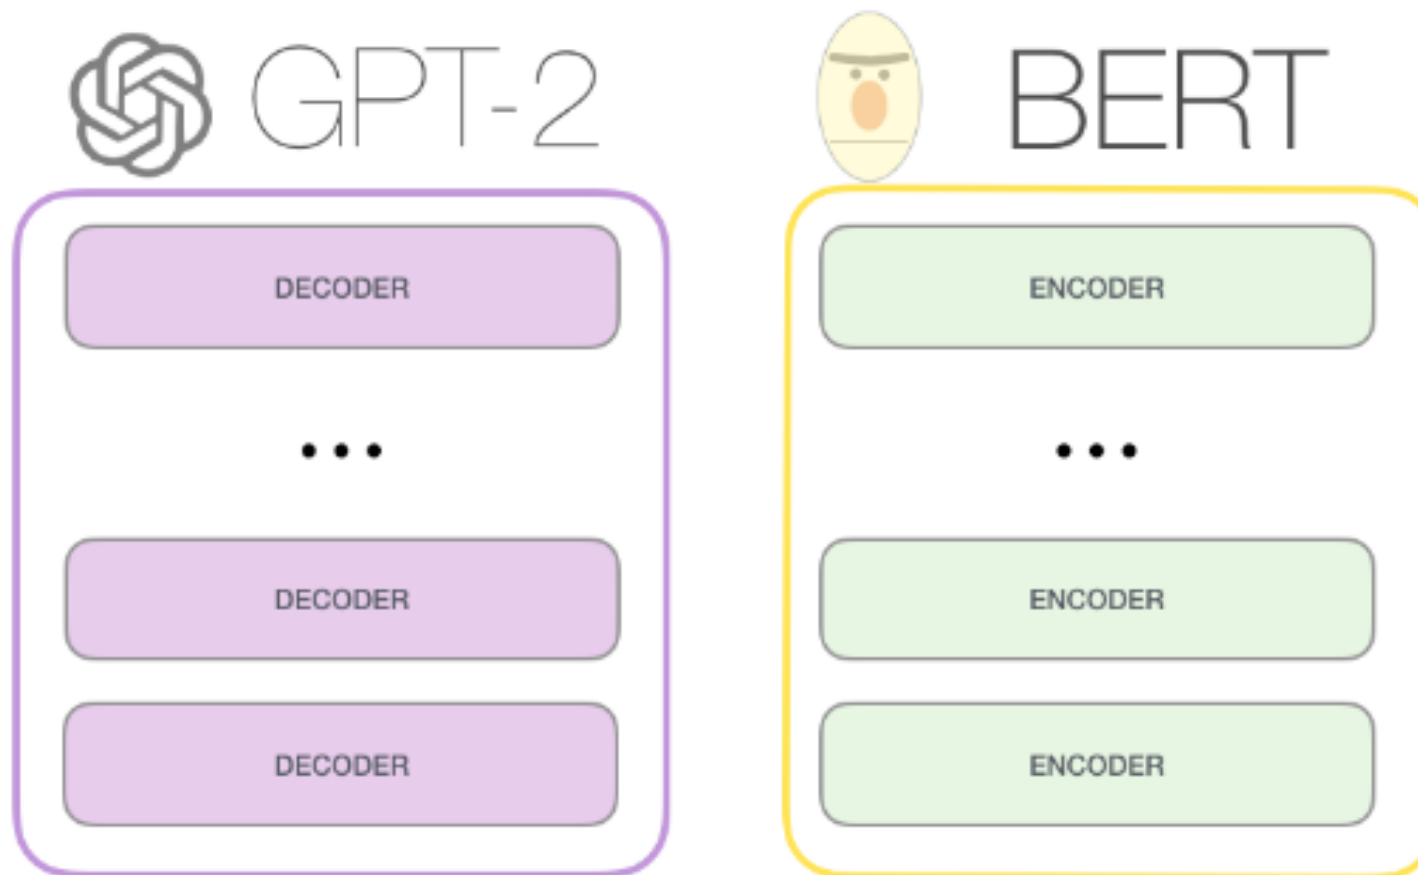
The encoder



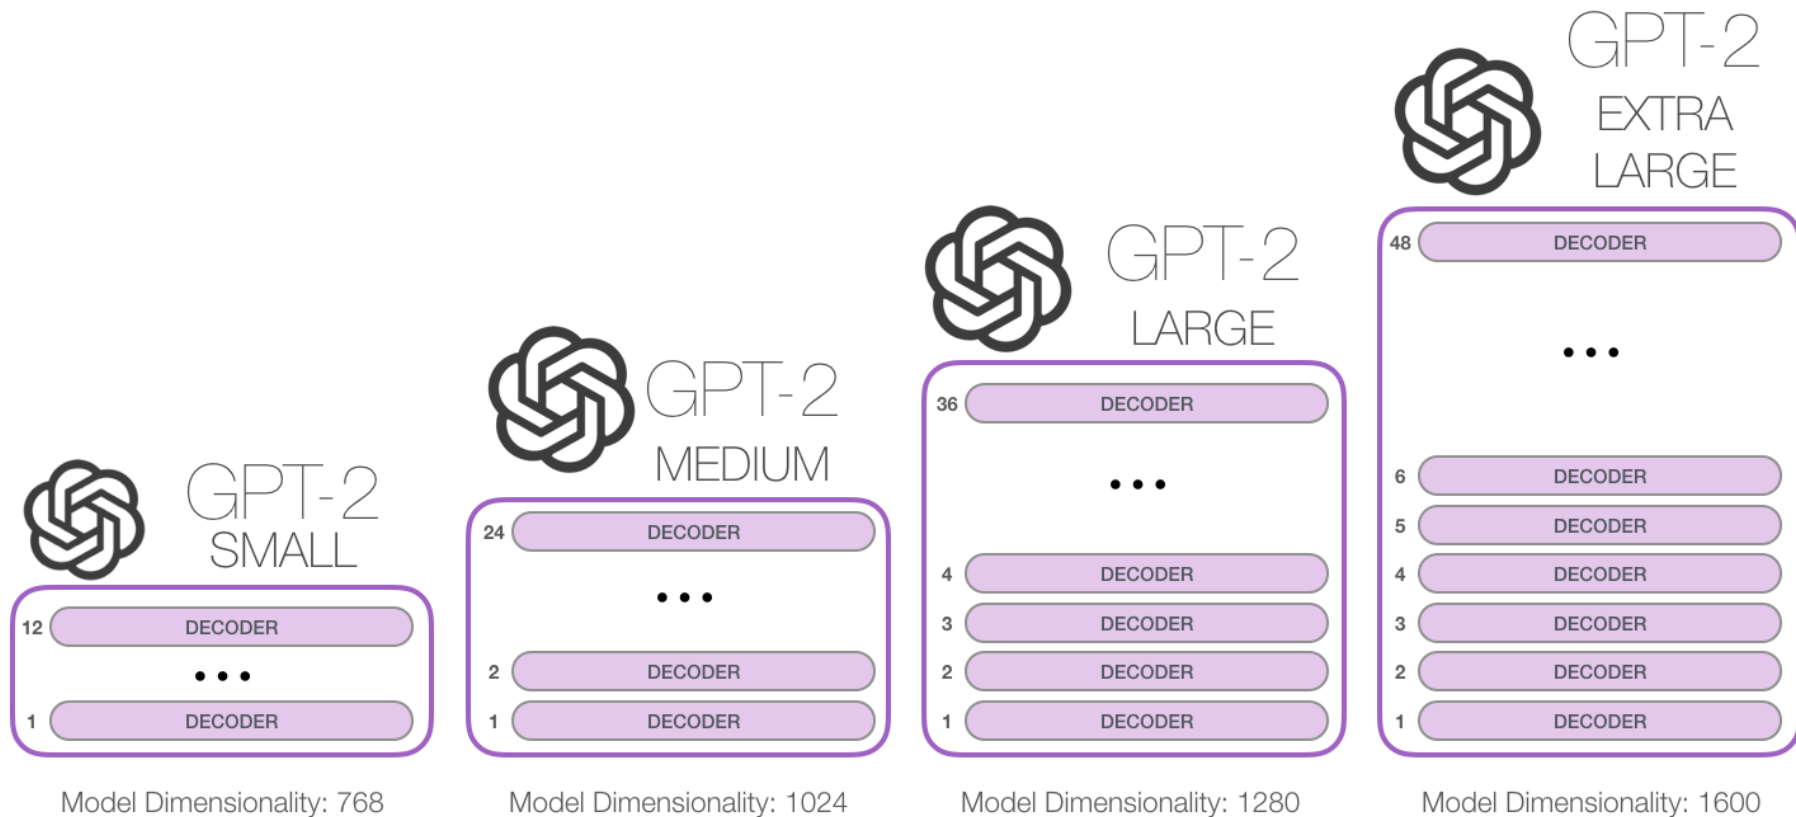
Transformers, GPT-2, and BERT

- ▶ A transformer uses an encoder stack to model the input, and a decoder stack to model the output.
- ▶ But if we don't have an input and we just want to model the "next word"
 - ▶ We can suppress the encoder side of a transformer and output the "next word" one by one
 - ▶ It gives us the GPT
- ▶ If we are only interested in forming a language model for input for other tasks
 - ▶ We don't need the transformer decoder,
 - ▶ It gives us BERT.

GPT 2 and BERT



GPT2



The embedding size varies according to the model (between 768 and 1600)

117 M
parameters


345 M
parameters

762 M
parameters

1,542 M
parameters

GPT2 in action

- Input: The GPT-2 can process 1024 tokens.
 - Each token flows through all the decoder blocks along its own path.
- Output: GPT2 produces one token at a time
 - This token is added to the input sequence to produce the following token



How to use GPT2

- ▶ The easiest way to use the transformers is through the API implemented by huggingface
 - ▶ To be preferred if you don't need to re-train the network.
 - ▶ huggingface.co
 - ▶ <https://github.com/huggingface/transformers>
- ▶ Use
 - ▶ Text generation
 - ▶ Text Summarization
- ▶ You can also install from scratch
 - ▶ Use if you need to specialize the network (fine tuning)
 - ▶ <https://medium.com/analytics-vidhya/gpt2-for-sentiment-analysis-38cd9832d5e9>

GPT2Tokenizer

```
>>> from transformers import GPT2Tokenizer
```

```
>>> tokenizer = GPT2Tokenizer.from_pretrained("gpt2")
```

```
>>> tokenizer("Hello world")['input_ids']  
[15496, 995]
```

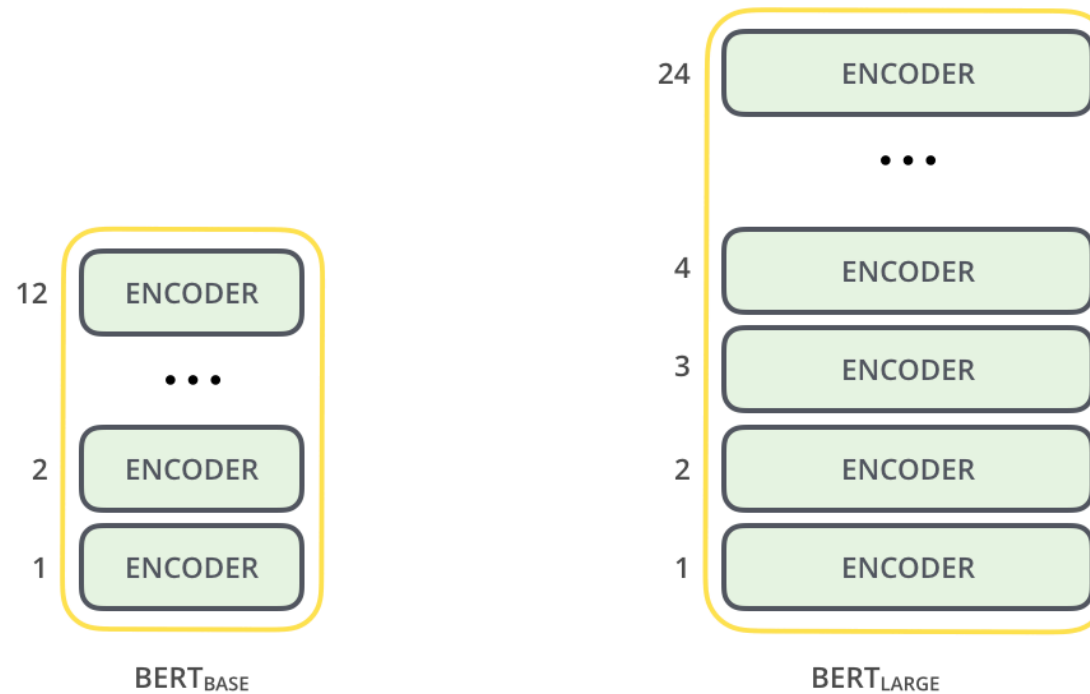
```
>>> tokenizer(" Hello world")['input_ids']  
[18435, 995]
```

- ▶ Remark: The tokenizer treats spaces as parts of a token.
 - ▶ A word will not be encoded identically depending on whether it is at the beginning of the sentence (without space) or not.

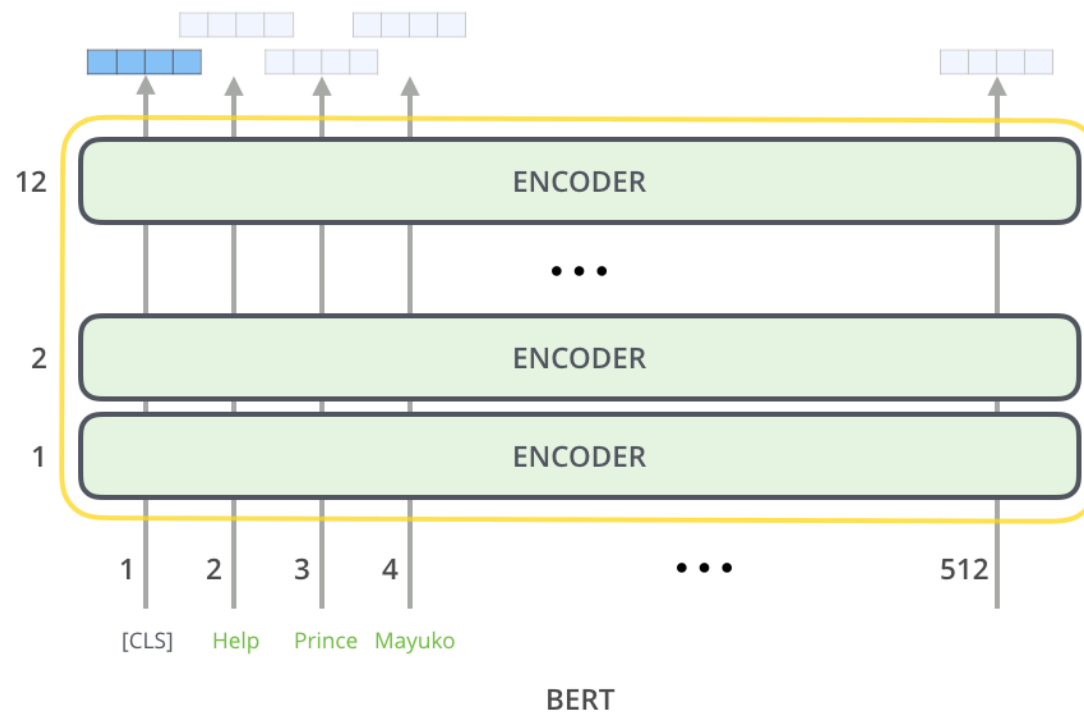
Text generation with GPT2...

- ▶ **Go to GPT2 notebook**
- ▶ **Need the installation of transformers library**
- ▶ `from transformers import GPT2Tokenizer, GPT2LMHeadModel`
- ▶ `tokenizer = GPT2Tokenizer.from_pretrained('distilgpt2')`
- ▶ `indexed_tokens = tokenizer.encode(text)`
- ▶ `model = GPT2LMHeadModel.from_pretrained('gpt2')`
- ▶ `outputs = model(tokens_tensor)`
- ▶ `predicted_index = torch.argmax(predictions[0, -1, :]).item()`
- ▶ `predicted_text = tokenizer.decode([predicted_index])`

Bidirectional Encoder Representation from Transformers (BERT)



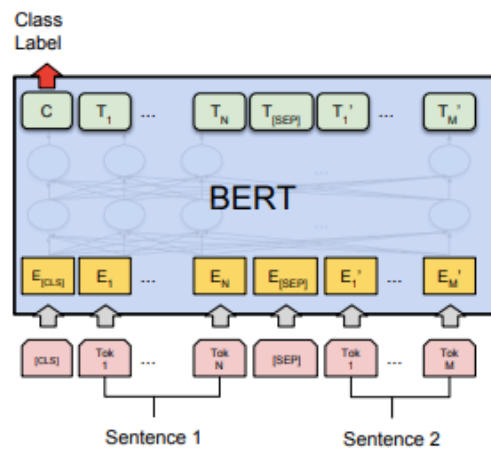
Model input and output



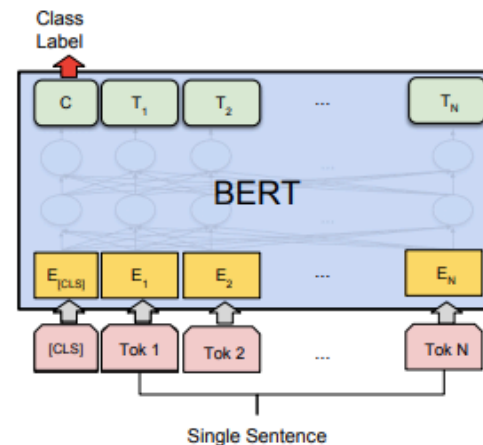
Input: 512 tokens

Output: 768 or 1024 regarding the model

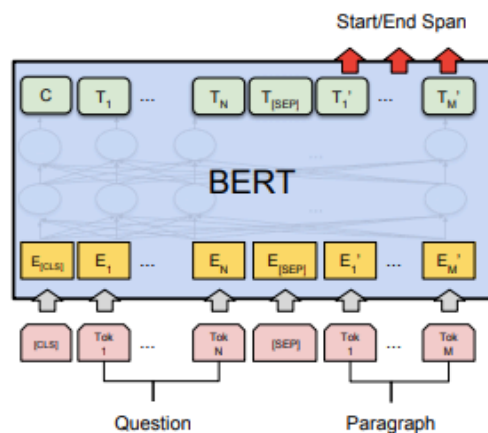
Main BERT usage with/without fine training



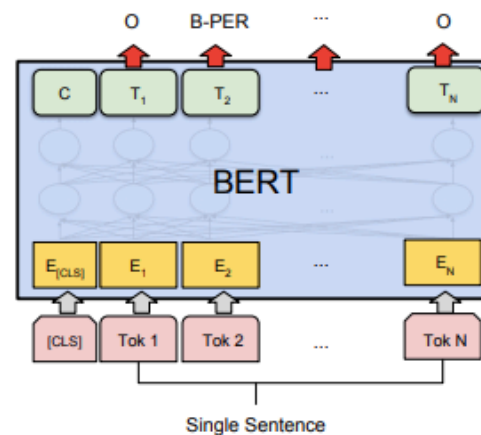
(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG



(b) Single Sentence Classification Tasks:
SST-2, CoLA



(c) Question Answering Tasks:
SQuAD v1.1



(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

How to use BERT

- ▶ Like GPT2, the easiest way to use the transformers is through the API implemented by huggingface
 - ▶ To be preferred if you don't need to re-train the network.
 - ▶ huggingface.co
 - ▶ <https://github.com/huggingface/transformers>
- ▶ Use
 - ▶ Text classification
 - ▶ NER task

BertTokenizer

```
>>> from transformers import BertTokenizer
```

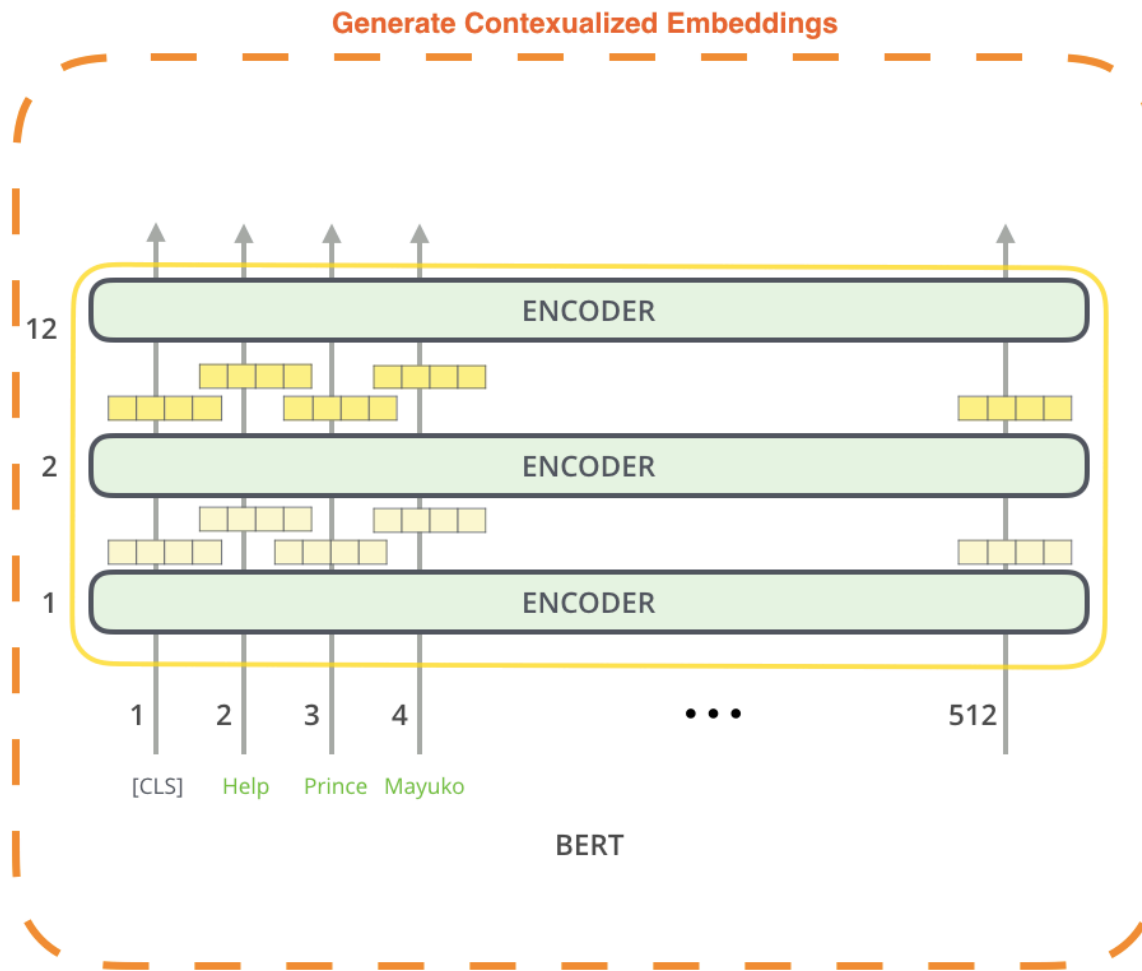
```
>>> tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
```

```
>>> tokenizer.tokenize("I take aspirin.")  
['i', 'take', 'as', '###pi', '###rin', '.']
```


```
>>> tokenizer.tokenize("I like chocolate")  
['i', 'like', 'chocolate']
```

- ▶ Remark: The tokenizer split OOV in sub-piece
 - ▶ the same token always has the same id
 - ▶ but the embedding changes

Feature extraction with BERT



The output of each encoder layer along each token's path can be used as a feature representing that token.



But which one should we use?

Bert embedding (feature extraction)

Use all encoder level

- ▶ `from transformers import BertTokenizer, TFBertModel`
- ▶ `tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')`
- ▶ `model = TFBertModel.from_pretrained('bert-base-uncased')`
- ▶ `tokenized_text = tokenizer.encode(review)`
- ▶ `input_ids = tf.constant(tokenized_text[:MAX_BERT_SIZE-2])`
- ▶ `outputs = model(input_ids[None, :])`
- ▶ `Prediction_scores, classification_scores = outputs[:2]`

`Prediction_scores` : embedding of each word


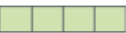






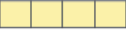
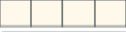





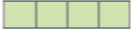


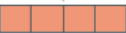



- ▶ Shape: nb_sentences, nb_tokens, 768

`Classification_scores` : last token (sentence embedding)

- ▶ Shape: nb_sentences, 768

Feature Extraction, which embedding to use?

What is the best contextualized embedding for “Help” in that context?
For named-entity recognition task CoNLL-2003 NER

		Dev F1 Score
12 	First Layer Embedding 	91.0
• • •		
7 	Last Hidden Layer 12 	94.9
6 		
5 	Sum All 12 Layers	95.5
4 	12  + • • • + 2  + 1  = 	
3 	Second-to-Last Hidden Layer 11 	95.6
2 		
1 	Sum Last Four Hidden	95.9
	12  + 11  + 10  + 9  = 	
Help	Concat Last Four Hidden	96.1
	9 10 11 12 	


Sentiment analysis with Bert

- ▶ **Go to Bert notebook**
- ▶ **Need the installation of transformers library**
- ▶ `from transformers import BertTokenizer`
- ▶ `BertTokenizer.from_pretrained("bert-base-uncased")`
- ▶ `marked_text = "[CLS] " + "I take aspirin. I like chocolate" + "[SEP]"`
- ▶ `tokenized_text = tokenizer.tokenize(marked_text)`
`['[CLS]', 'i', 'take', 'as', '##pi', '##rin', '.', 'i', 'like', 'chocolate', '[SEP]']`
- ▶ `encoded_text = tokenizer.encode(tokenized_text)`

Sentiment analysis with Bert

- ▶ `outputs = model(tf.constant(encoded_text))[None, :])`
- ▶ `prediction_scores, classification_scores = outputs[:2]`
- ▶ `prediction_scores.shape`
`(1, 10, 768)`
- ▶ `classification_scores.shape`
`(1, 768)`

Summary



Summary

- ▶ Transformer
 - ▶ Modelling language
 - ▶ Use the entire sentence before producing an output
 - ▶ The output could depend of the task
- ▶ LSTM is difficult to parallelize
- ▶ Self-Attention is a proposal to resolve the problems
- ▶ Transformer is sequence-to-sequence architecture
 - ▶ A set of encoders construct a latent representation of the input
 - ▶ A set of decoders could be use pour project the latent representation in.
a new space
- ▶ Specific lecture next year

Summary

- ▶ Transformer use very large model
 - ▶ Not so easy to use
 - ▶ Need computational resources
- ▶ 2 easy way to use transformer model
 - ▶ Tensorflow_hub library: <https://www.tensorflow.org/hub/installation>
 - ▶ Transformers library: <https://huggingface.co/transformers/>
- ▶ ELMO use character embedding and a bi-LSTM in order to produce an embedding based on word prediction
- ▶ Bert uses mainly the encoder part and could be used for
 - ▶ Word / Sentence embedding
 - ▶ Text classification or Sentiment Analysis
 - ▶ NER
 - ▶ Q&A or Text translation
- ▶ GPT uses mainly the decoder part and could be used for
 - ▶ Text summarization
 - ▶ Text generation