

Theory of Statistical Learning

Part II

Damien Garreau

Université Côte d'Azur

2021

Outline

1. Linear predictors

- Linear classification

- Linear regression

- Ridge regression

- Polynomial regression

- Logistic regression

2. Tree classifiers

- Partition rules

3. Boosting

1. Linear predictors

1.1. Linear classification

Linear functions

- ▶ $\mathcal{X} = \mathbb{R}^d$, $\mathcal{Y} = \mathbb{R}$
- ▶ thus $x_i = (x_{i,1}, x_{i,2}, \dots, x_{i,d})^\top$
- ▶ we consider no bias term (otherwise *affine*):

$$\{h : x \mapsto w^\top x, w \in \mathbb{R}^d\}.$$

- ▶ **Reminder:** given two vectors $u, v \in \mathbb{R}^d$,

$$\langle u, v \rangle = u^\top v = \sum_{j=1}^d u_j v_j.$$

- ▶ binary classification: 0-1 loss, $\mathcal{Y} = \{-1, +1\}$
- ▶ **Important:** compose h with $\phi : \mathbb{R} \rightarrow \mathcal{Y}$ (typically the sign)

Halfspaces

- ▶ thus our function class is

$$\mathcal{H} = \{x \mapsto \text{sign}(w^\top x), w \in \mathbb{R}^d\}.$$

- ▶ it is possible to show that $VC(\mathcal{H}) = d + 1$
- ▶ **Consequence:** \mathcal{H} is PAC learnable with sample complexity

$$\Omega\left(\frac{d + \log(1/\delta)}{\varepsilon}\right).$$

- ▶ **Important assumption:** data is linearly separable
- ▶ that is, there is a $w^* \in \mathbb{R}^d$ such that

$$y_i = \text{sign}(\langle w^*, x_i \rangle) \quad \forall 1 \leq i \leq n.$$

Linear programming

- ▶ **Empirical risk minimization:** recall that we are looking for w such that

$$\hat{\mathcal{R}}_S(w) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{y_i \neq \text{sign}(w^\top x_i)}$$

is minimal

- ▶ **Question:** how to solve this?
- ▶ we want $y_i = \text{sign}(w^\top x_i)$ for all $1 \leq i \leq n$
- ▶ equivalent formulation: $y_i \langle w, x_i \rangle > 0$
- ▶ we know that there is a vector that satisfies this condition (w^*)
- ▶ let us set $\gamma = \min_i \{y_i \langle w^*, x_i \rangle\}$ and $\bar{w} = w^* / \gamma$
- ▶ we have shown that there is a vector such that $y_i \langle \bar{w}, x_i \rangle \geq 1$ for any $1 \leq i \leq n$ (and it is an ERM)

Linear programming, ctd.

- ▶ define the matrix $A \in \mathbb{R}^{n \times d}$ such that

$$A_{i,j} = y_i x_{i,j}.$$

- ▶ **Intuition:** observations \times labels
- ▶ remember that we have the ± 1 label convention
- ▶ define $v = (1, \dots, 1)^\top \in \mathbb{R}^n$
- ▶ then we can rewrite the above problem as

$$\text{maximize } \langle u, w \rangle \text{ subject to } Aw \leq v,$$

with $u = 0$ for instance

- ▶ we call this sort of problems **linear programs**¹
- ▶ solvers readily available, e.g., `scipy.optimize.linprog` if you use Python

¹Boyd, Vandenberghe, *Convex optimization*, Cambridge University Press, 2004

The perceptron

- ▶ another possibility: the *perceptron*²
- ▶ **Idea:** iterative algorithm that constructs $w^{(1)}, w^{(2)}, \dots, w^{(T)}$
- ▶ update rule: at each step, find i that is misclassified and set

$$w^{(t+1)} = w^{(t)} + y_i x_i .$$

- ▶ **Question:** why does it work?
- ▶ pushes w in the right direction:

$$y_i \langle w^{(t+1)}, x_i \rangle = y_i \langle w^{(t)} + y_i x_i, x_i \rangle = y_i \langle w^{(t)}, x_i \rangle + \|x_i\|^2$$

- ▶ remember, we want $y_i \langle w, x_i \rangle > 0$ for all i

²Rosenblatt, *The perceptron, a perceiving and recognizing automaton*, tech report, 1957

Exercise

Exercise: Of course, one does not have to use the squared loss. Instead, we may prefer to use

$$\ell(y, y') = |y - y'|.$$

1. show that, for any $a \in \mathbb{R}$,

$$|c| = \min_{a \geq 0} a \quad \text{subject to} \quad c \leq a \text{ and } c \geq -a.$$

2. use the previous question to show that ERM with the absolute value loss function is equivalent to minimizing the linear function $\sum_{i=1}^n s_i$, where the s_i satisfy linear constraints
3. write it in matrix form, that is, find $A \in \mathbb{R}^{2n \times (n+d)}$, $v \in \mathbb{R}^{d+n}$, and $b \in \mathbb{R}^{2n}$ such that the LP can be written

$$\text{minimize } c^\top v \quad \text{subject to} \quad Av \leq b.$$

1.2. Linear regression

Least squares

- ▶ regression \Rightarrow squared-loss function

$$\ell(y, y') = (y - y')^2.$$

- ▶ still looking at linear functions:

$$\mathcal{H} = \{h : x \mapsto \langle w, x \rangle \text{ s.t. } w \in \mathbb{R}^d\}.$$

- ▶ empirical risk in this context:

$$\hat{\mathcal{R}}_S(h) = \frac{1}{n} \sum_{i=1}^n (w^\top x_i - y_i)^2 = F(w).$$

- ▶ also called **mean squared error**
- ▶ empirical risk minimization: we want to minimize $w \mapsto F(w)$ with respect to $w \in \mathbb{R}^d$
- ▶ F is a **convex, smooth** function

Least squares, ctd.

- let us compute the gradient of F :

$$\begin{aligned}\frac{\partial F}{\partial w_j}(w) &= \frac{1}{n} \sum_{i=1}^n \frac{\partial}{\partial w_j} (w^\top x_i - y_i)^2 \\ &= \frac{1}{n} \sum_{i=1}^n 2 \frac{\partial}{\partial w_j} w^\top x_i (w^\top x_i - y_i) \\ \frac{\partial F}{\partial w_j}(w) &= \frac{2}{n} \sum_{i=1}^n (w^\top x_i - y_i) x_{i,j} .\end{aligned}$$

Least squares, ctd.

- ▶ we can rewrite it, define

$$A = \sum_{i=1}^n x_i x_i^\top \text{ and } b = \sum_{i=1}^n y_i x_i ,$$

then solving $\nabla F(w) = 0$ is equivalent to

$$Aw = b .$$

- ▶ if A is invertible, straightforward:

$$\hat{w} = A^{-1}b$$

- ▶ what happens when A is not invertible?

Singular value decomposition

- ▶ since A is symmetric, it has an eigendecomposition

$$A = VDV^{\top},$$

with $D \in \mathbb{R}^d$ diagonal and V orthonormal

- ▶ define D^+ such that

$$D_{i,i}^+ = 0 \text{ if } D_{i,i} = 0 \text{ and } D_{i,i}^+ = \frac{1}{D_{i,i}} \text{ otherwise.}$$

- ▶ define $A^+ = VD^+V^{\top}$
- ▶ then we set

$$\hat{w} = A^+ b.$$

Singular value decomposition, ctd.

- ▶ why did we do that?
- ▶ let v_i denote the i th column of V , then

$$A\hat{w} = AA^+b \quad (\text{definition of } \hat{w})$$

$$= VDV^T VD^+ V^T b \quad (\text{definition of } A^+)$$

$$= VDD^+ V^T b \quad (V \text{ is orthonormal})$$

$$A\hat{w} = \sum_{i:D_{i,i} \neq 0} v_i v_i^T b.$$

- ▶ in definitive, $A\hat{w}$ is the projection of b onto the span of v_i such that $D_{i,i} \neq 0$
- ▶ since the span of these v_i is the span of the x_i and b is in the linear span of the x_i , we have $A\hat{w} = b$

Recap

- ▶ **What happens when we invoke**
`sklearn.linear_model.LinearRegression` with default parameters?
- ▶ `fit_intercept` is `True` → assumes that the data is not centered (our maths are not totally accurate)
- ▶ `normalize` is `False` → we are responsible for the normalization of our data
- ▶ behind the scenes, calls `scipy.linalg.lstsq` when fitting, which itself calls LAPACK (Linear Algebra PACKage)³
- ▶ LAPACK is coded in Fortran90



³<http://www.netlib.org/lapack/>

1.3. Ridge regression

Ridge regression

- ▶ same hypothesis class: linear functions

$$\mathcal{H} = \{h : x \mapsto w^\top x, w \in \mathbb{R}^d\}$$

- ▶ squared loss:

$$\ell(y, y') = (y - y')^2.$$

- ▶ **Idea:** regularization:

$$\text{minimize } \left\{ \frac{1}{n} \sum_{i=1}^n (y_i - w^\top x_i)^2 + \lambda \|w\|^2 \right\},$$

with $\|u\|^2 = u_1^2 + \dots + u_d^2$ and $\lambda > 0$ a *regularization parameter*

Exercise

Exercise: Let $(x_1, y_1), \dots, (x_n, y_n) \in \mathbb{R}^d \times \mathbb{R}$ be n given training samples. For any $w \in \mathbb{R}^d$, set

$$F(w) = \frac{1}{n} \sum_{i=1}^n (y_i - w^\top x_i)^2 + \lambda \|w\|^2.$$

Notice that F is a convex smooth function and find its minimizer \hat{w} in closed-form. Recall that we defined

$$A = \sum_{i=1}^n x_i x_i^\top \quad \text{and} \quad b = \sum_{i=1}^n y_i x_i.$$

Recap

- ▶ **What happens when we invoke** `sklearn.linear_model.Ridge` with default settings?
- ▶ $\alpha = 1 \rightarrow \lambda = 1/n$ with our notation, barely any regularization if n large
- ▶ `fit_intercept` is `True` \rightarrow does not consider centered data (so our analysis is not entirely accurate)
- ▶ `normalize` is `False` \rightarrow we decide whether we normalize our data
- ▶ `solver` is `auto` \rightarrow `sklearn` will decide how to solve the minimization problem depending on the size of the data: **the solution could be not exact!**
- ▶ `tol` = 0.001 \rightarrow tolerance threshold on the residuals

1.4. Polynomial regression

Polynomial regression

- ▶ linear regression is a powerful tool, especially because we can **transform the inputs**

- ▶ **Example:** polynomial regression in \mathbb{R}

- ▶ inputs $x_1, \dots, x_n \in \mathbb{R}$

- ▶ define the mapping $\phi(x) = (1, x, x^2, \dots, x^p)^\top$

- ▶ then

$$\langle w, \phi(x) \rangle = w_0 + w_1x + w_2x^2 + \dots + w_px^p,$$

and we can find the best coefficients by linear regression

- ▶ `numpy.polyfit` → very handy when we want to fit univariate data

1.5. Logistic regression

Logistic regression

- ▶ regression task, but the output is $\mathcal{Y} = [0, 1]$: we predict **the probability of belonging to class 1**
- ▶ hypothesis class:

$$\mathcal{H} = \{x \mapsto \phi(\langle w, x \rangle), w \in \mathbb{R}^d\},$$

with ϕ the *logistic function* (aka *sigmoid function*)

$$\phi(z) = \frac{1}{1 + e^{-z}}.$$

- ▶ logistic loss:

$$\ell(y, y') = \frac{-1}{2}(1 + y) \log y' + \frac{1}{2}(1 - y) \log(1 - y').$$

Logistic function




Figure: the logistic function $\sigma : t \mapsto 1/(1 + e^{-t})$.

Exercise

Exercise: show that empirical risk minimization with the logistic loss is equivalent to minimizing

$$F(w) = \sum_{i=1}^n \log(1 + \exp(-y_i \langle w, x_i \rangle)).$$

Is F a convex function of w ? Compute the gradient of F with respect to w . Can you solve $\nabla F(w) = 0$?

Recap

- ▶ **What happens when we call**
`sklearn.linear_model.LogisticRegression?`
- ▶ penalty is $\ell_2 \rightarrow$ **there is regularization by default!** (not much though, $C = 1$)
- ▶ `fit_intercept` is `True` \rightarrow again, our maths are not entirely accurate
- ▶ `solver` is `liblinear` \rightarrow since there is no closed-form, a solver will be used
- ▶ `liblinear` uses coordinate descent
- ▶ will default soon to `lbfgs` (Limited-memory Broyden-Fletcher-Goldfarb-Shanno)
- ▶ **do not worry too much about the solvers**, just change if you see that it is not converging

2. Tree classifiers

2.1. Partition rules

Introduction

- ▶ let $\mathcal{X} = \mathbb{R}^d$ and $\mathcal{Y} = \mathbb{R}$
- ▶ in this section, we consider partition-based classifiers:

$$\mathcal{H} = \left\{ h : x \mapsto \sum_{j=1}^p h_j \mathbb{1}_{x \in A_j} \right\},$$

where $a_j \in \mathbb{R}$ and A_1, \dots, A_p form a *partition* of the space

- ▶ that is,

$$A_1 \cup \dots \cup A_p = \mathcal{X} \quad \text{and} \quad A_i \cap A_j = \emptyset \forall i \neq j.$$

- ▶ the A_j s are often called *cells*
- ▶ generally, for practical reasons the A_j s are rectangles

ERM for partition rules

- ▶ assume that the partition is fixed
- ▶ regression with squared loss, then the ERM rule gives

$$h_j = \frac{1}{|\{i \text{ s.t. } i \in A_j\}|} \sum_{i \in A_j} x_i,$$

that is the average of the observations on each cell

- ▶ classification \Rightarrow majority vote
- ▶ thus **ERM** \Leftrightarrow finding the best partition (for a fixed p)
- ▶ **Problem:** this is computationally very hard! p^n possibilities to compare
- ▶ even if we restrict ourselves to rectangles, intractable

Trees

- ▶ one possible solution: start from \mathcal{X} and *split* iteratively
- ▶ we obtain a tree-like structure
- ▶ another advantage in doing so: root the new data efficiently
- ▶ **Question:** how do we make the splits?
- ▶ **general answer:** take an heuristic that makes sense
- ▶ each heuristic yields a different algorithm, completed with stopping criterion (do a split only if gain greater than γ)
- ▶ **Notation:** I current node, I_L (resp. I_R) left (resp. right) node after the split
- ▶ **Note:** we focus on classification ($\mathcal{Y} = \{0, 1\}$)

ID3⁵ and C4.5

Definition: Let S be a finite set of points. Then we define the *entropy* of S by

$$H(S) = \sum_{y \in \mathcal{Y}} -p(y) \log_2 p(y),$$

where $p(y)$ is the proportion of elements of S classified as y .

- ▶ easy to see that $H(S) = 0$ means that S is perfectly classified ($0 \log 0 = 0$)
- ▶ C4.5 criterion:⁴ find direction and split that maximizes

$$H(I) - H(I_L) - H(I_R).$$

⁴Quinlan, *C4.5: Programs for Machine Learning*, 1993

⁵Quinlan, *Induction of decision trees*, Machine Learning, 1986

CART

- ▶ later supplanted by CART trees⁶

Definition: Let S be a finite set of points. We define the *Gini impurity* by

$$G(S) = \sum_{y \in \mathcal{Y}} p(y)(1 - p(y)).$$

- ▶ CART trees: find direction and split that minimizes

$$G(I) - G(I_L) - G(I_R).$$

- ▶ for regression, variance reduction criterion

⁶Breiman et al., *Classification and Regression Trees*, 1984

When to stop?

- ▶ usually, many direction to try: CART reduces to a random subset of directions
- ▶ also possible to specify T a max height for the tree
- ▶ other strategy: grow the trees to the full extent, and then **pruning**

Recap

- ▶ **What happens by default when we invoke the function `sklearn.tree.DecisionTreeClassifier`?** let us look at least at the main options
- ▶ `criterion` is set to Gini → we are using CART trees
- ▶ `splitter` is set to best → looking at the best split at each step
- ▶ `max_depth` is None → splitting until leaves are pure or contain less than `min_samples_split`
- ▶ `min_samples_split` = 2
- ▶ `max_features` is None → no max number of features, log could be a reasonable choice if we have many features
- ▶ `max_leaf_nodes`: None → many leaves, we could also restrict this
- ▶ `min_impurity_decrease` = 0 → continues to split even if very small gain

3. Boosting

Introduction

- ▶ **Idea:** aggregate many weak classifiers together, then majority voting
- ▶ **Examples:** linear classifier, trees,...