

Theory of Statistical Learning

Part II

Damien Garreau

Université Côte d'Azur

2021

Outline

1. Linear predictors

- Linear classification

- Linear regression

- Ridge regression

- Polynomial regression

- Logistic regression

2. Tree-based classifiers

- Partition rules

- Random forests

3. Boosting

- Adaboost

- Gradient boosting

- XGBoost

1. Linear predictors

1.1. Linear classification

Linear functions

- ▶ $\mathcal{X} = \mathbb{R}^d$, $\mathcal{Y} = \mathbb{R}$
- ▶ thus $x_i = (x_{i,1}, x_{i,2}, \dots, x_{i,d})^\top$
- ▶ we consider no bias term (otherwise *affine*):

$$\{h : x \mapsto w^\top x, w \in \mathbb{R}^d\}.$$

- ▶ **Reminder:** given two vectors $u, v \in \mathbb{R}^d$,

$$\langle u, v \rangle = u^\top v = \sum_{j=1}^d u_j v_j.$$

- ▶ binary classification: 0-1 loss, $\mathcal{Y} = \{-1, +1\}$
- ▶ **Important:** compose h with $\phi : \mathbb{R} \rightarrow \mathcal{Y}$ (typically the sign)

The sign function

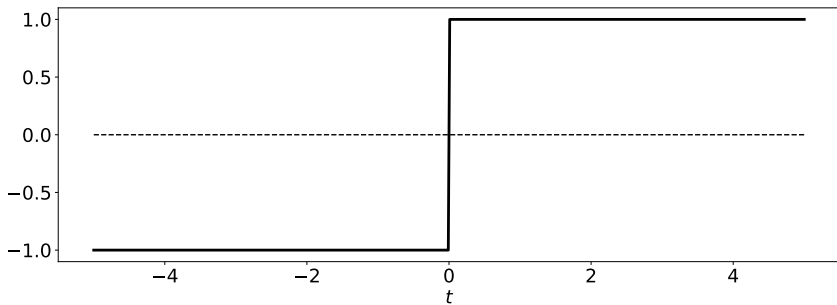


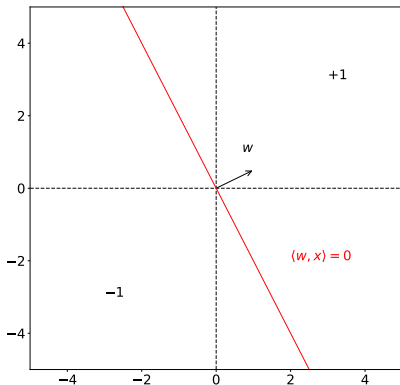
Figure: the sign function

Halfspaces

- ▶ thus our function class is

$$\mathcal{H} = \{x \mapsto \text{sign}(w^\top x), w \in \mathbb{R}^d\}.$$

- ▶ gives label +1 to vector pointing in the same direction as w



VC dimension of halfspaces

Proposition: the VC dimension of halfspaces in dimension d is exactly $d + 1$.

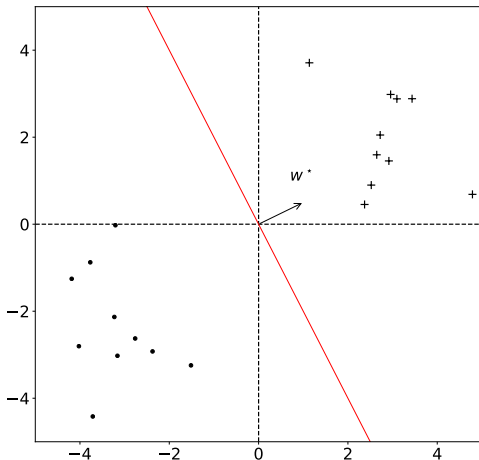
► **Consequence:** \mathcal{H} is PAC learnable with sample complexity

$$\Omega\left(\frac{d + \log(1/\delta)}{\varepsilon}\right).$$

Linearly separable data

- ▶ **Important assumption:** data is linearly separable
- ▶ that is, there is a $w^* \in \mathbb{R}^d$ such that

$$y_i = \text{sign}(\langle w^*, x_i \rangle) \quad \forall 1 \leq i \leq n.$$



Linear programming

- ▶ **Empirical risk minimization:** recall that we are looking for w such that

$$\hat{\mathcal{R}}_S(w) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{y_i \neq \text{sign}(w^\top x_i)}$$

is minimal

- ▶ **Question:** how to solve this?
- ▶ we want $y_i = \text{sign}(w^\top x_i)$ for all $1 \leq i \leq n$
- ▶ equivalent formulation: $y_i \langle w, x_i \rangle > 0$
- ▶ we know that there is a vector that satisfies this condition (w^*)
- ▶ let us set $\gamma = \min_i \{y_i \langle w^*, x_i \rangle\}$ and $\bar{w} = w^* / \gamma$
- ▶ we have shown that there is a vector such that $y_i \langle \bar{w}, x_i \rangle \geq 1$ for any $1 \leq i \leq n$ (and it is an ERM)

Linear programming, ctd.

- ▶ define the matrix $A \in \mathbb{R}^{n \times d}$ such that

$$A_{i,j} = y_i x_{i,j}.$$

- ▶ **Intuition:** observations \times labels
- ▶ remember that we have the ± 1 label convention
- ▶ define $v = (1, \dots, 1)^\top \in \mathbb{R}^n$
- ▶ then we can rewrite the above problem as

$$\text{maximize } \langle u, w \rangle \text{ subject to } Aw \leq v,$$

with $u = 0$ for instance

- ▶ we call this sort of problems **linear programs**¹
- ▶ solvers readily available, e.g., `scipy.optimize.linprog` if you use Python

¹Boyd, Vandenberghe, *Convex optimization*, Cambridge University Press, 2004

The perceptron

- ▶ another possibility: the *perceptron*²
- ▶ **Idea:** iterative algorithm that constructs $w^{(1)}, w^{(2)}, \dots, w^{(T)}$
- ▶ update rule: at each step, find i that is misclassified and set

$$w^{(t+1)} = w^{(t)} + y_i x_i .$$

- ▶ **Question:** why does it work?
- ▶ pushes w in the right direction:

$$y_i \langle w^{(t+1)}, x_i \rangle = y_i \langle w^{(t)} + y_i x_i, x_i \rangle = y_i \langle w^{(t)}, x_i \rangle + \|x_i\|^2$$

- ▶ remember, we want $y_i \langle w, x_i \rangle > 0$ for all i

²Rosenblatt, *The perceptron, a perceiving and recognizing automaton*, tech report, 1957

1.2. Linear regression

Least squares

- ▶ regression \Rightarrow squared-loss function

$$\ell(y, y') = (y - y')^2.$$

- ▶ still looking at linear functions:

$$\mathcal{H} = \{h : x \mapsto \langle w, x \rangle \text{ s.t. } w \in \mathbb{R}^d\}.$$

- ▶ empirical risk in this context:

$$\hat{\mathcal{R}}_S(h) = \frac{1}{n} \sum_{i=1}^n (w^\top x_i - y_i)^2 = F(w).$$

- ▶ also called **mean squared error**
- ▶ empirical risk minimization: we want to minimize $w \mapsto F(w)$ with respect to $w \in \mathbb{R}^d$
- ▶ F is a **convex, smooth** function

Least squares, ctd.

- ▶ let us compute the gradient of F :

$$\begin{aligned}\frac{\partial F}{\partial w_j}(w) &= \frac{1}{n} \sum_{i=1}^n \frac{\partial}{\partial w_j} (w^\top x_i - y_i)^2 \\&= \frac{1}{n} \sum_{i=1}^n 2 \cdot \frac{\partial}{\partial w_j} (w^\top x_i - y_i) \cdot (w^\top x_i - y_i) \\&= \frac{1}{n} \sum_{i=1}^n 2 \cdot \frac{\partial}{\partial w_j} (\cdots + w_j x_{i,j} + \cdots - y_i) \cdot (w^\top x_i - y_i) \\ \frac{\partial F}{\partial w_j}(w) &= \frac{2}{n} \sum_{i=1}^n x_{i,j} \cdot (w^\top x_i - y_i).\end{aligned}$$

Least squares, ctd.

- ▶ it is more convenient to write $\nabla F(w) = 0$ in matrix notation
- ▶ define $X \in \mathbb{R}^{n \times d}$ the matrix such that line i of X is observation x_i
- ▶ one can check that, for any $1 \leq j, k \leq d$,

$$(X^\top X)_{j,k} = \sum_{i=1}^n x_{i,j} x_{i,k} .$$

- ▶ thus

$$\begin{aligned}(X^\top X w)_j &= \sum_{k=1}^d (X^\top X)_{j,k} w_k \\ &= \sum_{k=1}^d \sum_{i=1}^n x_{i,j} x_{i,k} w_k \\ &= \sum_{i=1}^n x_{i,j} w^\top x_i .\end{aligned}$$

Least squares, ctd.

- ▶ thus, if we define

$$A = X^T X = \sum_{i=1}^n x_i x_i^T \in \mathbb{R}^{d \times d} \text{ and } b = X^T y = \sum_{i=1}^n y_i x_i \in \mathbb{R},$$

solving $\nabla F(w) = 0$ is equivalent to solving

$$Aw = b.$$

- ▶ if A is invertible, straightforward:

$$\hat{w} = A^{-1}b$$

- ▶ computational cost: $\mathcal{O}(d^3)$ (inversion is actually a bit less)
- ▶ what happens when A is not invertible?

Singular value decomposition

- ▶ since A is symmetric, it has an eigendecomposition

$$A = VDV^{\top},$$

with $D \in \mathbb{R}^d$ diagonal and V orthonormal

- ▶ define D^+ such that

$$D_{i,i}^+ = 0 \text{ if } D_{i,i} = 0 \text{ and } D_{i,i}^+ = \frac{1}{D_{i,i}} \text{ otherwise.}$$

- ▶ define $A^+ = VD^+V^{\top}$
- ▶ then we set

$$\hat{w} = A^+ b.$$

Singular value decomposition, ctd.

- ▶ why did we do that?
- ▶ let v_i denote the i th column of V , then

$$A\hat{w} = AA^+b \quad (\text{definition of } \hat{w})$$

$$= VDV^{\top}VD^+V^{\top}b \quad (\text{definition of } A^+)$$

$$= VDD^+V^{\top}b \quad (V \text{ is orthonormal})$$

$$A\hat{w} = \sum_{i:D_{i,i} \neq 0} v_i v_i^{\top} b.$$

- ▶ in definitive, $A\hat{w}$ is the projection of b onto the span of v_i such that $D_{i,i} \neq 0$
- ▶ since the span of these v_i is the span of the x_i and b is in the linear span of the x_i , we have $A\hat{w} = b$
- ▶ cost of SVD: $\mathcal{O}(dn^2)$ if $d > n$ (SVD of X)

Exercise

Exercise: Of course, one does not have to use the squared loss. Instead, we may prefer to use

$$\ell(y, y') = |y - y'|.$$

1. show that, for any $c \in \mathbb{R}$,

$$|c| = \min_{a \geq 0} a \quad \text{subject to} \quad a \geq c \text{ and } a \geq -c.$$

2. use the previous question to show that ERM with the absolute value loss function is equivalent to minimizing the linear function $\sum_{i=1}^n s_i$, where the s_i satisfy linear constraints
3. write it in matrix form, that is, find $A \in \mathbb{R}^{2n \times (n+d)}$, $v \in \mathbb{R}^{d+n}$, and $b \in \mathbb{R}^{2n}$ such that the LP can be written

$$\text{minimize } c^\top v \quad \text{subject to} \quad Av \leq b.$$

Correction of the exercise

1. The absolute value is the smallest positive number larger than both c and $-c$ for any real number c .
2. In that case, the empirical risk can be written

$$\hat{\mathcal{R}}_S(w) = \frac{1}{n} \sum_{i=1}^n |y_i - w^\top x_i|.$$

We deduce the result from question 1.

3. One possibility is to define $v = (w_1, \dots, w_d, s_1, \dots, s_n)^\top \in \mathbb{R}^{n+d}$, $c = (0, \dots, 0, 1, \dots, 1)^\top \in \mathbb{R}^{d+n}$, $b = (y_1, \dots, y_n, -y_1, \dots, -y_n)^\top \in \mathbb{R}^{2n}$, and

$$A = \begin{pmatrix} -X & -I_n \\ X & -I_n \end{pmatrix} \in \mathbb{R}^{2n \times (n+d)},$$

with $X \in \mathbb{R}^{n \times d}$ the matrix whose lines are the x_i s and I_n the identity matrix.

Recap

- ▶ **What happens when we invoke**
`sklearn.linear_model.LinearRegression` with default parameters?
- ▶ `fit_intercept` is `True` → assumes that the data is not centered (our maths are not totally accurate)
- ▶ `normalize` is `False` → we are responsible for the normalization of our data
- ▶ behind the scenes, calls `scipy.linalg.lstsq` when fitting, which itself calls LAPACK (Linear Algebra PACKage)³
- ▶ LAPACK is coded in Fortran90



³<http://www.netlib.org/lapack/>

1.3. Ridge regression

Ridge regression

- ▶ same hypothesis class: linear functions

$$\mathcal{H} = \{h : x \mapsto w^\top x, w \in \mathbb{R}^d\}$$

- ▶ squared loss:

$$\ell(y, y') = (y - y')^2.$$

- ▶ **Idea:** regularization:

$$\text{minimize } \left\{ \frac{1}{n} \sum_{i=1}^n (y_i - w^\top x_i)^2 + \lambda \|w\|^2 \right\},$$

with $\|u\|^2 = u_1^2 + \dots + u_d^2$ and $\lambda > 0$ a *regularization parameter*

Exercise

Exercise: Let $(x_1, y_1), \dots, (x_n, y_n) \in \mathbb{R}^d \times \mathbb{R}$ be n given training samples. For any $w \in \mathbb{R}^d$, set

$$F(w) = \frac{1}{n} \sum_{i=1}^n (y_i - w^\top x_i)^2 + \lambda \|w\|^2.$$

Notice that F is a convex smooth function.

1. show that the minimizer \hat{w} satisfies

$$(X^\top X + n\lambda I_d) w = X^\top y.$$

2. show that $X^\top X + n\lambda I_d$ is an invertible matrix

Correction of the exercise

1. Let $1 \leq j \leq d$ and let us compute $\partial_j F$:

$$\begin{aligned}\frac{\partial F}{\partial w_j}(w) &= \frac{\partial}{\partial w_j} \left(\frac{1}{n} \sum_{i=1}^n (y_i - w^\top x_i)^2 \right) + \frac{\partial}{\partial w_j} (\lambda(w_1^2 + \dots + w_d^2)) \\ &= \frac{2}{n} \sum_{i=1}^n x_{i,j} \cdot (w^\top x_i - y_i) + 2\lambda w_j,\end{aligned}$$

where we used the derivation for the least squares. We deduce the result by setting to zero and multiplying by n .

Correction of the exercise, ctd.

2. By contradiction, suppose that $X^\top X + n\lambda I_d$ is not invertible. Then

$$\det(X^\top X + n\lambda I_d) = 0.$$

In other words, $-n\lambda$ is an eigenvalue of $X^\top X$. Since $X^\top X$ is a symmetric matrix, its spectrum is $\subseteq \mathbb{R}$. Moreover, it is positive definite, thus all eigenvalues are non-negative. Since $\lambda > 0$, we deduce that $-n\lambda$ cannot be an eigenvalue of $X^\top X$ and we can conclude. \square

Recap

- ▶ **What happens when we invoke** `sklearn.linear_model.Ridge` with default settings?
- ▶ $\alpha = 1 \rightarrow \lambda = 1/n$ with our notation, barely any regularization if n large
- ▶ `fit_intercept` is `True` \rightarrow does not consider centered data (so our analysis is not entirely accurate)
- ▶ `normalize` is `False` \rightarrow we decide whether we normalize our data
- ▶ `solver` is `auto` \rightarrow `sklearn` will decide how to solve the minimization problem depending on the size of the data: **the solution could be not exact!**
- ▶ `tol` = 0.001 \rightarrow tolerance threshold on the residuals

1.4. Polynomial regression

Polynomial regression

- ▶ linear regression is a powerful tool, especially because we can transform the inputs in a non-linear fashion

- ▶ **Example:** polynomial regression in \mathbb{R}

- ▶ inputs $x_1, \dots, x_n \in \mathbb{R}$

- ▶ define the mapping $\phi(x) = (1, x, x^2, \dots, x^p)^\top$

- ▶ then

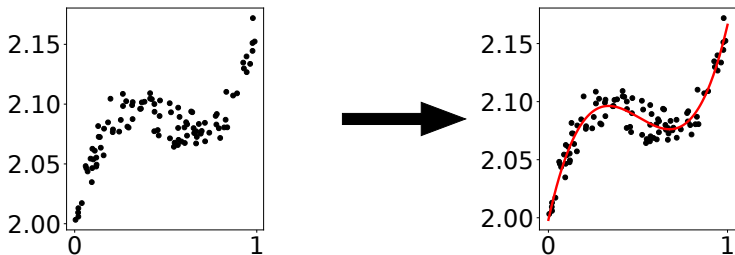
$$\langle w, \phi(x) \rangle = w_0 + w_1x + w_2x^2 + \dots + w_px^p,$$

and we can find the best coefficients by linear regression

- ▶ `numpy.polyfit` → very handy when we want to fit univariate data

Polynomial regression, ctd.

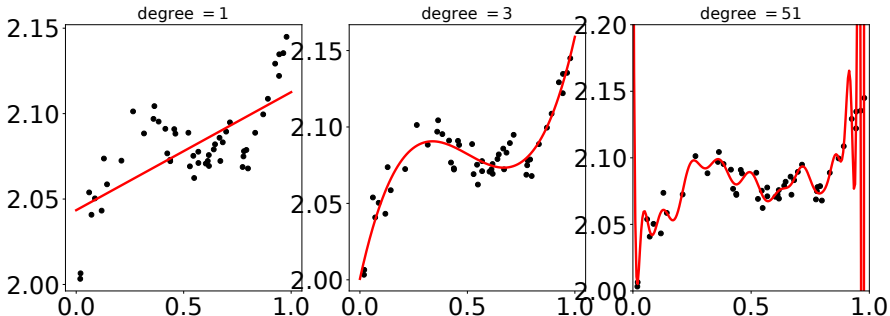
- ▶ **Example:** data = degree three polynomial + Gaussian noise with small variance
- ▶ fit a degree 3 polynomial:



- ▶ **Remark:** in practice, we do not know the degree of the polynomial!

Polynomial regression, ctd.

- ▶ typical case of under / overfitting:
 - ▶ when degree too low, poor fit
 - ▶ when degree too high, wiggly function ($n + 1 \Rightarrow$ interpolation)



1.5. Logistic regression

Logistic regression

- ▶ classification with $\mathcal{Y} = \{0, 1\}$
- ▶ however, we predict **the probability of belonging to class 1**
- ▶ hypothesis class:

$$\mathcal{H} = \{x \mapsto \phi(\langle w, x \rangle), w \in \mathbb{R}^d\},$$

with ϕ the *logistic function* (aka *sigmoid function*)

$$\phi(z) = \frac{1}{1 + e^{-z}}.$$

- ▶ **Intuition:** squeeze the score between 0 and 1 to transform it into a probability
- ▶ $\mathbb{P}(y = 1 | x) = \phi(w^\top x)$ and $\mathbb{P}(y = 0 | x) = 1 - \phi(w^\top x)$

Logistic function

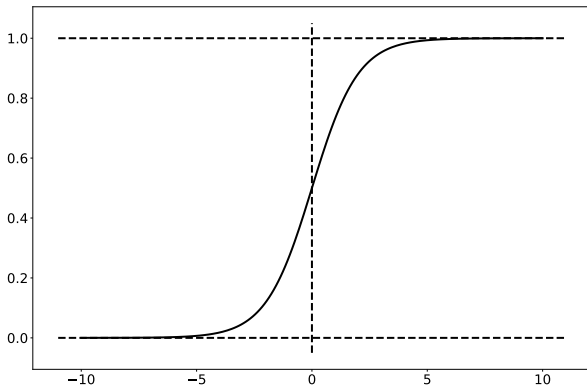
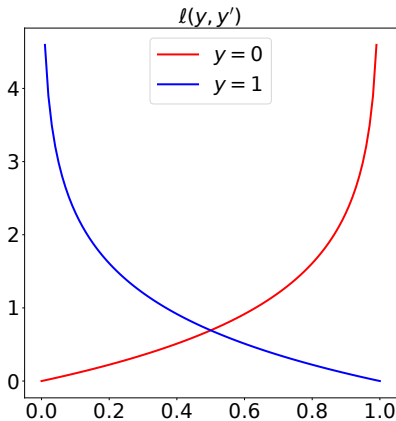


Figure: the logistic function $\phi : t \mapsto 1/(1 + e^{-t})$.

Logistic loss

- ▶ **Next:** we need to define a loss function
- ▶ for any y, y' , we define the *logistic loss*:

$$\ell(y, y') = -(1 - y) \log(1 - y') - y \log y'.$$



Logistic regression


- ▶ finally, logistic regression = empirical risk minimization with the logistic loss
- ▶ that is, minimize for $w \in \mathbb{R}^d$

$$\hat{\mathcal{R}}_S(w) = \sum_{i=1}^n \{-(1 - y_i) \log(1 - \phi(w^\top x_i)) - y_i \log \phi(w^\top x_i)\} .$$

- ▶ **Remark (i):** we can show that this is equivalent to maximum likelihood for a certain prior distribution
- ▶ **Remark (ii):** complicated to optimize (see exercise)

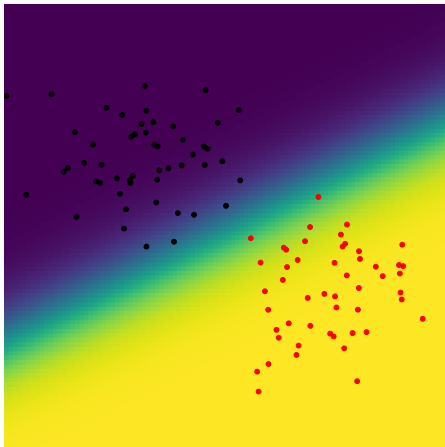
Logistic regression in dimension 1

► **Example:** in dimension one:



Logistic regression in dimension 2

- **Example:** in dimension two:



Exercise

Exercise: Recall that we defined the logistic loss by

$$\ell(y, y') = -(1 - y) \log(1 - y') - y \log y'.$$

1. Show that ERM with the logistic loss is equivalent to minimizing

$$F(w) = \sum_{i=1}^n \log(1 + \exp(-\tilde{y}_i \langle w, x_i \rangle)),$$

where $\tilde{y}_i = \text{sign}(y_i - 0.5)$. Deduce that $\hat{\mathcal{R}}$ is a convex function of w .

2. Compute the gradient of $\hat{\mathcal{R}}$ with respect to w . *Hint:* show that $\phi'(z) = \phi(z)(1 - \phi(z))$.
3. Can you solve $\nabla \hat{\mathcal{R}}(w) = 0$? If not, propose a strategy for finding a good w .