

Decision Trees and Random Forests

Taki Eddine MEKHALFA
Taki-Eddine.MEKHALFA@inria.fr

Simple strategy:

1. **Split** (stratify, segment) the input space into disjoint regions

Simple strategy:

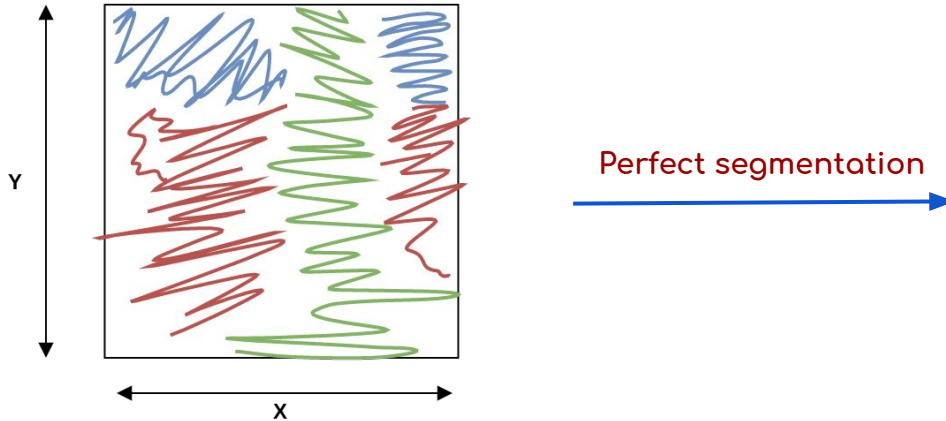
1. **Split** (stratify, segment) the input space into disjoint regions
2. For each input, **find the region** in to which it belongs

Simple strategy:

1. **Split** (stratify, segment) the input space into disjoint regions
2. For each input, **find the region** in to which it belongs
3. Return the **prediction associated with this region** (mean or mode for example)


Simple strategy:

1. **Split** (stratify, segment) the input space into disjoint regions
2. For each input, **find the region** in to which it belongs
3. Return the **prediction associated with this region** (mean or mode for example)




Simple strategy:

1. **Split** (stratify, segment) the input space into disjoint regions
2. For each input, **find the region** in to which it belongs
3. Return the **prediction associated with this region** (mean or mode for example)



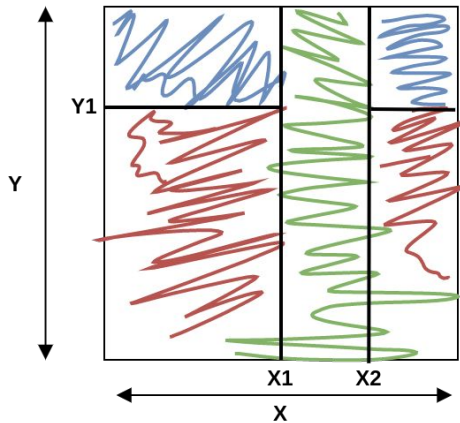
Simple strategy:

1. **Split** (stratify, segment) the input space into disjoint regions
2. For each input, **find the region** in to which it belongs
3. Return the **prediction associated with this region** (mean or mode for example)





Simple strategy:


1. **Split** (stratify, segment) the input space into disjoint regions
2. For each input, **find the region** in to which it belongs
3. Return the **prediction associated with this region** (mean or mode for example)







Equivalent Decision Tree

















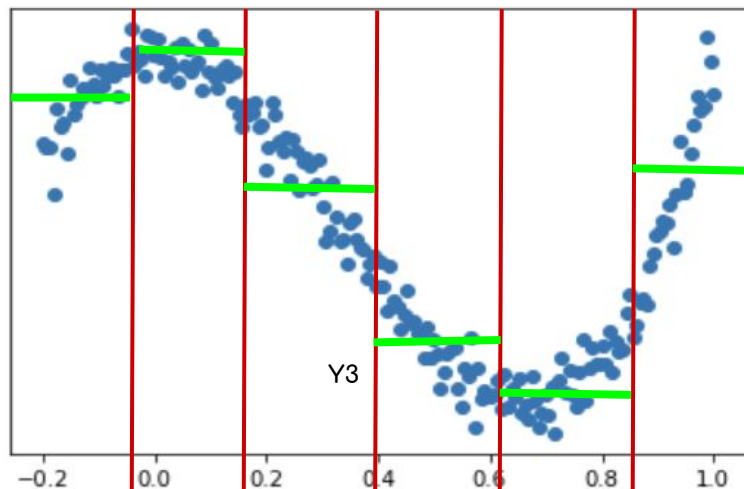













1. Same thing for **high dimensional data** (you can segment your space using different **features**)
2. Same thing for **classification** (use **majority class** instead of the **mean**)




- Decision trees are generally built only to a maximum depth and not until there is only one data point inside leaf nodes. We prefer moderately deep trees than very deep trees. Why?


- Decision trees are generally built only to a maximum depth and not until there is only one data point inside leaf nodes. We prefer moderately deep trees than very deep trees. Why?
 - We don't build very deep trees to avoid overfitting

- Decision trees are generally built only to a maximum depth and not until there is only one data point inside leaf nodes. We prefer moderately deep trees than very deep trees. Why?
 - We don't build very deep trees to avoid overfitting
- We have a limited number of splits. What features we choose to split on? And in what order? How to split a given feature?

Split on some feature: What boundary?




Tree 1




Tree 2

Split on some feature: What boundary?




Tree 1




Tree 2

Cost(Tree) = $\sum \text{Cost}(\text{Region}^i)$ s.t. $\text{Cost}(\text{Region}^i) = \text{MSE}(\text{Region}^i) = \text{MSE}(y^j - m^i)$ where m^i is the mean of Region^i

Split on some feature: What boundary?



Tree 1

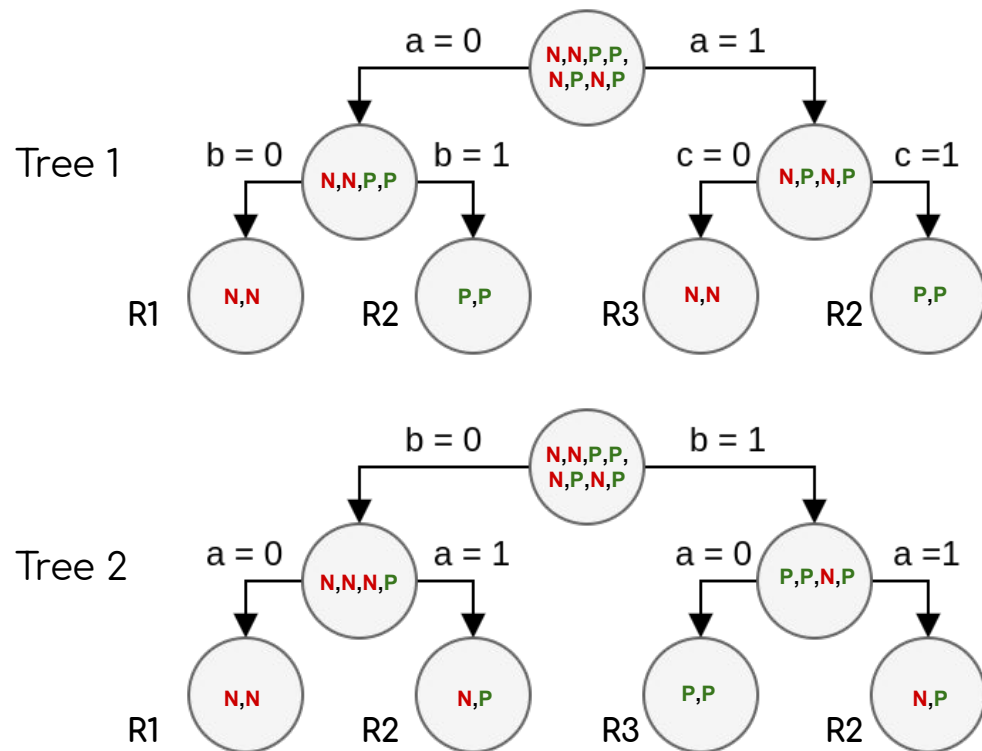


Tree 2

Cost(Tree) = $\sum \text{Cost}(\text{Region}^i)$ s.t. $\text{Cost}(\text{Region}^i) = \text{MSE}(\text{Region}^i) = \text{MSE}(y^j - m^i)$ where m^i is the mean of Region^i

Multiple features: In what order?

a	b	c	CLASS
0	0	0	NEG
0	0	1	NEG
0	1	0	POS
0	1	1	POS
1	0	0	NEG
1	0	1	POS
1	1	0	NEG
1	1	1	POS



$$\text{Cost}(\text{Region}^i) = \text{Impurity}(\text{Region}^i)$$


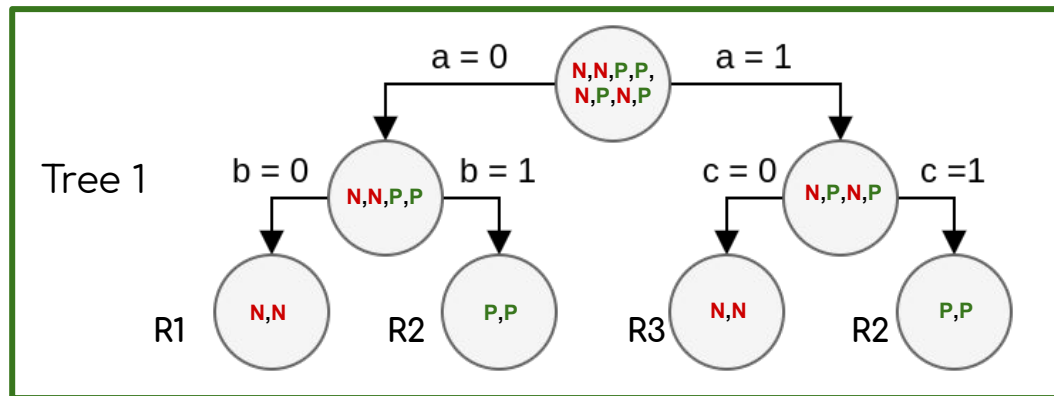
Multiple features: In what order?

a	b	c	CLASS
0	0	0	NEG
0	0	1	NEG
0	1	0	POS
0	1	1	POS
1	0	0	NEG
1	0	1	POS
1	1	0	NEG
1	1	1	POS

Annotations for feature splits:

- $a = 0$ (split on feature a)
- $b = 0$ (split on feature b)
- $b = 1$ (split on feature b)
- $c = 0$ (split on feature c)
- $c = 1$ (split on feature c)
- $a = 1$ (split on feature a)

$$\text{Cost}(\text{Region}^i) = \text{Impurity}(\text{Region}^i)$$




$$\text{Cost}(\text{Region}^i) = \text{Impurity}(\text{Region}^i)$$

P,P,P,P Is purer than P,P,P,N Is purer than P,P,N,N


How to quantify Impurity?

First Idea: 1 - percentage of the majority class



← The more the majority class is major the smaller the impurity

Second Idea: Entropy

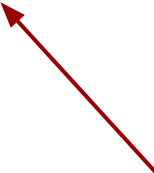


← The more the classes are equally probable, the more entropy (impurity) is high and vice versa

$$\text{Entropy} = -\sum \log(\text{percentage}(C^j)) \times \text{percentage}(C^j)$$

- ❑ Inputs:
 - ❑ Depth of the tree: K
 - ❑ Learning data points
- A. Algorithm:
 - a. For all trees of depth K choose the best tree with respect to the defined loss

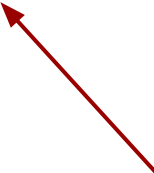
- ❑ Inputs:
 - ❑ Depth of the tree: K
 - ❑ Learning data points
- A. Algorithm:
 - a. For all trees of depth K choose the best tree with respect to the defined loss



Non feasible (actually NP-Complete).
There is an intractable number of **K -depth** trees

- ❑ Inputs:
 - ❑ Depth of the tree: K
 - ❑ Learning data points
- A. Algorithm:
 - a. For all trees of depth K choose the best tree with respect to the defined loss

What can we do?

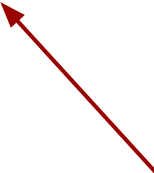


Non feasible (actually NP-Complete).
There is an intractable number of **K-depth** trees

- ❑ Inputs:
 - ❑ Depth of the tree: K
 - ❑ Learning data points
- A. Algorithm:
 - a. For all trees of depth K choose the best tree with respect to the defined loss

What can we do?

Optimize greedily!




Non feasible (actually NP-Complete).
There is an intractable number of K -depth trees

- Inputs:

- Depth of the tree: K
 - Learning data points

- A. Algorithm:

- a. For each internal node with depth = $d < K$:
 - i. Choose the feature f s.t when splitted it decreases the loss the most.
 - ii. Split on f and create two binary nodes and increment their depth to $d + 1$





Feasible

Greedy algorithm:

1. Simple and efficient
2. **BUT** we are not guaranteed to build the optimal tree
3. **Nonetheless** works well in practice.

Greedy algorithm:


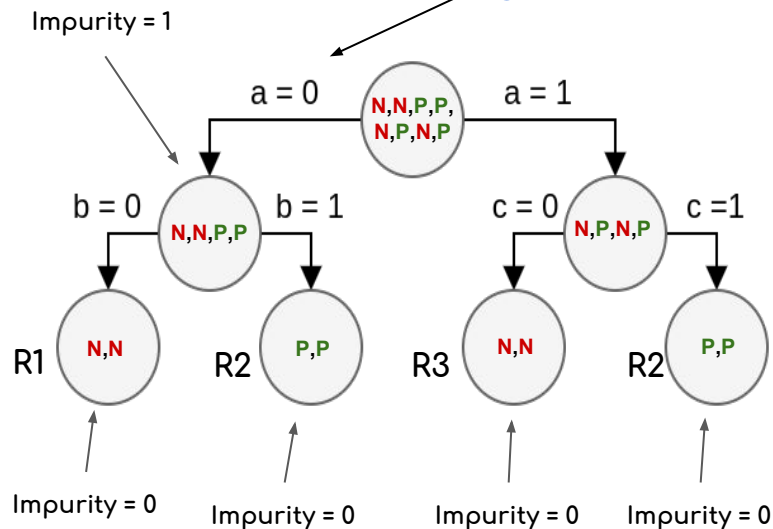
1. Simple and efficient
2. **BUT** we are not guaranteed to build the optimal tree
3. **Nonetheless** works well in practise.



Greedy algorithm:

1. Simple and efficient
2. **BUT** we are not guaranteed to build the optimal tree
3. **Nonetheless** works well in practise.

Seems worthless locally, but globally optimal



Take away points:

- Decision trees are **simple**
- Decision trees are **interpretable** (their behavior **can be explained** using splitting rules) and mimic rule based human decision making (trees can be easily visualized)
- Decision trees are built using a **greedy** algorithm which usually **works well in practice**.
- Decision trees are a flexible model: **Small bias** (fits non linear relationships) but a **big variance** (a small change in data might result in a large change of the final tree i.e **overfit easily**)

Decision trees are overfit easily (high variance)! How to remediate this problem?

Decision trees are overfit easily (high variance)! How to remediate this problem?

Suppose you have n independent variables $\{Z_1, Z_2, \dots, Z_n\}$ where $\text{var}(Z_i) = \sigma^2$,

then $\text{var}(\overline{Z_n}) = \sigma^2/n$ i.e averaging reduces the variance.

Decision trees are overfit easily (high variance)! How to remediate this problem?

Suppose you have n independent variables $\{Z_1, Z_2, \dots, Z_n\}$ where $\text{var}(Z_i) = \sigma^2$,

then $\text{var}(\overline{Z_n}) = \sigma^2/n$ i.e averaging reduces the variance.

Idea => **Construct B decision trees** on **B independent data sets** and **average their results** at prediction time (or use majority vote in case of classification for example)

Decision trees are overfit easily (high variance)! How to remediate this problem?

Suppose you have n independent variables $\{Z_1, Z_2, \dots, Z_n\}$ where $\text{var}(Z_i) = \sigma^2$,

then $\text{var}(\overline{Z_n}) = \sigma^2/n$ i.e averaging reduces the variance.

Idea => **Construct B decision trees** on **B independent data sets** and **average their results** at prediction time (or use majority vote in case of classification for example)

Problem => **We don't have B data sets!**

Decision trees are overfit easily (high variance)! How to remediate this problem?

Suppose you have n independent variables $\{Z_1, Z_2, \dots, Z_n\}$ where $\text{var}(Z_i) = \sigma^2$,

then $\text{var}(\overline{Z_n}) = \sigma^2/n$ i.e averaging reduces the variance.

Idea => **Construct B decision trees** on **B independent data sets** and **average their results** at prediction time (or use majority vote in case of classification for example)

Problem => **We don't have B data sets!**

Solution => **Use bootstrapping!**

Decision trees are overfit easily (high variance)! How to remediate this problem?

Suppose you have n independent variables $\{Z_1, Z_2, \dots, Z_n\}$ where $\text{var}(Z_i) = \sigma^2$,

then $\text{var}(\overline{Z_n}) = \sigma^2/n$ i.e averaging reduces the variance.

Idea => **Construct B decision trees** on **B independent data sets** and **average their results** at prediction time (or use majority vote in case of classification for example)

Problem => **We don't have B data sets!**


Solution => **Use bootstrapping!**

Bootstrapping: A sampling technique where for each data set with n points, sample n points **uniformly with replacement** (i.e a given data point might be sampled multiple times)

Theoretically, Bagging (Bootstrap aggregating), can be used on any prediction model. We show here a pseudo algorithm on bagging with regression decision trees:

- ❑ Inputs:
 - ❑ Number of trees B
 - ❑ Depth of a tree
- A. Algorithm:
 - a. For b in range(B):
 - i. Bootstrap a data set D^b of size n
 - ii. Fit a decision tree T^b on D^b
 - iii. Add T^b to the already built trees
- B. Prediction:
 - a. Take a input i
 - b. For each b compute $T^b(i)$
 - c. Return the average of all predictions

Theoretically, Bagging (Bootstrap aggregating), can be used on any prediction model. We show here a pseudo algorithm on bagging with regression decision trees:



Out-of-Bag error estimation:

=> Because of **bootstrapping**, each tree is constructed using **only a subset of the total data points** (one can show that a tree is constructed with **only ~ 2/3 of data**)

Out-of-Bag error estimation:

=> Because of **bootstrapping**, each tree is constructed using **only a subset of the total data points** (one can show that a tree is constructed with **only ~ 2/3 of data**)

=> random **1/3 of data** is not used to construct a given tree i.e each data point is not used in **1/3 of trees**

Out-of-Bag error estimation:

=> Because of **bootstrapping**, each tree is constructed using **only a subset of the total data points** (one can show that a tree is constructed with **only ~ 2/3 of data**)

=> random **1/3 of data** is not used to construct each tree i.e each data point is not used in **1/3 of trees**

=> This remaining data can be used as a validation subset to approximate test errors
=> No need for cross validation

Take away points:

- Increasing the total number of trees **will not cause overfitting**, therefore one can increase the number of trees until the error settles

Take away points:

- Increasing the total number of trees **will not cause overfitting**, therefore one can increase the number of trees until the error settles
- Individual trees can be built to deep levels, **we don't care about individual overfitting** since **this effect is overridden by averaging**

Take away points:

- Increasing the total number of trees **will not cause overfitting**, therefore one can increase the number of trees until the error settles
- Individual trees can be built to deep levels, **we don't care about individual overfitting** since **this effect is overridden by averaging**
- If there are strong features, **all trees will use them in first levels** => Trees will have similar behavior (**correlated**) => We don't gain much by averaging => **Use Random Forests**

How can we **decorrelate** different trees ?

How can we **decorrelate** different trees ?

=> for each tree, **use a subset of features only** (not all features)

How can we **decorrelate** different trees ?

=> for each tree, **use a subset of features only** (not all features)

=> Generally we build a tree using **m** randomly chosen features (where **m** < **p** the total number of features)

=> In practice we use **m** = \sqrt{p}

How can we **decorrelate** different trees ?

=> for each tree, **use a subset of features only** (not all features)

=> Generally we build a tree using **m** randomly chosen features (where **m** < **p** the total number of features)

=> In practice we use **m** = \sqrt{p}

=> When stronger features are not chosen, other features has more chances to be explored => **trees are different and uncorrelated**

Take away points:

- Random forests is just a bagging of decision trees but trees are built on a **subset of features only**.
- When $m = p$, random forests = bagging